NeedleFinder 2026: Sub-Millimeter Catheter Reconstruction on MRI for Interstitial Gynecologic Brachytherapy in Seconds

October 2026 (Target Journal: Physics in Medicine & Biology*)*

ABSTRACT

Objective. In MRI-guided interstitial gynecologic brachytherapy, a clinical team inserts 15 to 30 thin plastic catheters into a pelvic tumor and acquires a 3T T2-weighted MRI scan while the patient is under anesthesia. Before radiation can be planned, a physicist must trace the 3D path of every catheter on the scan. Doing this by hand takes 30 to 45 minutes. Existing automatic methods based on voxel segmentation struggle on T2-weighted MRI because plastic catheters appear only as tiny dark dots that look much like blood vessels, and brief signal dropouts break single catheters into disconnected pieces. Here we present NeedleFinder 2026, an algorithm that reconstructs the entire catheter bundle automatically, to sub-millimeter shaft accuracy, in a few seconds.

Approach. NeedleFinder 2026 combines imaging physics with three stages of machine learning. First, an analytical 3D filter (Stage 0) locates the entry plane of the perineal template automatically and measures circular dark-core symmetry on each slice—highlighting round catheter cross-sections while ignoring one-sided tissue borders. Second, a native-resolution 3D residual U-Net (Stage 1) predicts the sub-pixel center and local 3D tilt of each catheter on every slice. To prevent this network from mistaking pelvic veins for catheters, we pre-train it on 51 external needle-free pelvic MRI scans injected with synthetic catheter voids and fine-tune it on consensus traces from two expert teams. Third, a graph-based tracking pipeline (Stage 2) links detections across slices, stitches fragments broken by signal dropouts, evaluates the entire 15-to-30-catheter bundle together using a graph transformer, and pinpoints each catheter's deepest tip using a 1D axial neural network equipped with calibrated uncertainty intervals. We evaluated the pipeline across 54 clinical 3T MRI scans (1,063 catheters) using strictly nested 5 × 5 patient-level cross-validation.

Main results. NeedleFinder 2026 detected 91.4% of all 1,063 catheters with 86.9%–87.4% precision and an F1 score of 89.1%–89.4% (and 90.0%–90.5% F1 on the 759 full-depth treatment shafts), outperforming five peer-reviewed MRI/curvilinear methods and one recent tracking preprint reproduced on the identical nested splits by 6.0 to 50.7 F1 points ($p \leq 1.1 \times 10^{-5}$). Along the catheter shaft, the median in-plane error was 0.47 mm (mean 0.58–0.59 mm; 0.44 mm median on full-depth shafts). At the deepest tip, the median 3D error was 2.26–2.27 mm (mean 6.23 mm), with 98.8% of the residual tip variance lying longitudinally along the catheter track. Across 40 scans traced independently by two human expert teams, three-cornered-hat analysis and one-sided non-inferiority tests showed that the algorithm is

statistically non-inferior to human inter-observer agreement on both shafts ($p = 1.3 \times 10^{-8}$) and tips ($p = 2.6 \times 10^{-7}$). Post-implant CT scans from 38 paired patients confirmed that 52.0% of the algorithm's apparent "false alarms" on MRI were real physical catheters missed by the primary human team on MRI. In high-dose-rate ^{192}Ir radiation treatment simulations, reviewing just the 2 to 5 most uncertain tips flagged by the algorithm's conformal intervals took under a minute of simulated review and kept the median tumor dose error ($|\Delta D_{90}|$) to 2.8% (mean 5.9%). Without retraining, the gynecologic pipeline generalized directly to 3T prostate brachytherapy MRI (96.3%–100.0% F1), balanced steady-state MRI (92.9% F1), live 4D needle-insertion scans (82.8% F1), and post-implant CT. End-to-end reconstruction takes 4.15 s per patient on a workstation CPU (1.26 s on standard-resolution scans) and 2.67 s on a GPU.

Significance. Combining a simple physical symmetry prior with multi-stage graph learning turns a 30-to-45-minute manual tracing bottleneck into a 4-second automatic reconstruction followed by a brief, uncertainty-guided check in 3D Slicer.

PLAIN-LANGUAGE SUMMARY

Advanced cancers of the cervix, uterus, and vagina are often treated with *interstitial brachytherapy*. While the patient is asleep under anesthesia, doctors guide 15 to 30 thin plastic tubes (catheters) into the tumor. A tiny radioactive source will later travel through these tubes for a few minutes to destroy the cancer from the inside while sparing nearby healthy organs like the bladder and bowel.

Before any radiation can be delivered, the medical team takes an MRI scan to see where the tubes landed inside the tumor. Because radiation strength drops off very rapidly with distance, being off by even a couple of millimeters can under-treat the tumor or over-irradiate a healthy organ. Currently, a physicist must sit at a computer while the patient remains under anesthesia and trace every single tube by hand across dozens of MRI slices—a painstaking job that takes 30 to 45 minutes.

Automating this step on MRI has been difficult for a simple physical reason: plastic tubes contain no mobile water, so they do not glow on an MRI scan. Instead, each tube appears only as a tiny dark dot, two millimeters wide, on each image slice. The human pelvis is full of other dark structures—blood vessels, air pockets, and tissue borders—so standard artificial intelligence programs easily confuse blood vessels with tubes, or break a single tube into pieces when the dark dot fades on one slice.

We built **NeedleFinder 2026** to solve this problem in the way a careful human physicist does—by combining local image appearance with the overall geometry of the implant:

1. **Look for dark circles, not dark edges (Stage 0).** A simple mathematical filter checks whether each dark spot is surrounded by brighter tissue in *all* directions (like a tube cross-section) rather than dark on one side only (like a muscle border). It also automatically finds where the needle-guiding template sits at the base of the scan.

2. **Pinpoint tube centers and directions in 3D (Stage 1).** A 3D neural network—first pre-trained on 51 external pelvic MRI scans to tell artificial tubes apart from real blood vessels, and then trained on traces from two independent teams of human experts—marks the exact center and tilt of each tube on every slice, even when two tubes touch.
3. **Connect the dots into a coherent bundle and find each tip (Stage 2).** Two graph networks connect the marks from slice to slice, bridge gaps where a tube briefly fades, and judge whether all 15 to 30 curves make geometric sense together as a single implant. Finally, a specialized 1D network looks along the end of each tube to pinpoint its deepest tip and provides an uncertainty bar telling the physicist which tips are clear and which ones deserve a quick second look.

Tested across 54 patients (1,063 tubes), NeedleFinder 2026 found 91.4% of all tubes automatically. Along the length of each tube, its path agreed with human experts to within half a millimeter (0.47 mm median)—closer than two teams of human experts agreed with each other (0.49 mm median). At the deepest tip, its typical error was 2.27 mm, compared with 3.24 mm between the two human teams. Even more strikingly, when we checked post-implant CT scans (where tubes are easy to see), more than half (52%) of the "extra" tubes found by our software on MRI turned out to be real tubes that the primary human team had missed on MRI.

In practice, NeedleFinder 2026 runs inside the free **3D Slicer** medical software in **4 seconds** on a standard workstation computer (1.3 seconds on standard-resolution scans). Instead of tracing 25 tubes from scratch, the physicist clicks one button, checks the 2 to 5 tube tips that the software flags as uncertain in under a minute of simulated review, and proceeds to radiation planning with tumor dose accuracy within 2.8% (median) of a fully manual plan.

---

1 INTRODUCTION

1.1 *The Clinical Bottleneck in MRI-Guided Brachytherapy*

For locally advanced cervical and vaginal cancer, high-dose-rate (HDR) brachytherapy—delivering radiation directly inside the tumor—is essential for cure [1,2]. The international EMBRACE-I study demonstrated 92% five-year local tumor control by delivering a median dose of 90 Gy (EQD2) to the high-risk tumor volume [2], and the EMBRACE-II protocol prescribes 90 to 95 Gy to the target while strictly limiting dose to the adjacent bladder, rectum, and sigmoid colon [1]. When a tumor is large, asymmetric, or extends into surrounding tissues, a central uterine tandem alone cannot cover it safely. Clinicians therefore perform *interstitial* brachytherapy: alongside the central applicator, they insert 10 to 30 thin plastic catheters through a perineal template or freehand directly into the tumor [3–6].

T2-weighted (T2W) MRI is the gold standard for visualizing gynecologic tumors and surrounding soft-tissue organs, and international guidelines (GEC-ESTRO

and AAPM Task Group 303) recommend MRI-based treatment planning [7–10]. Ideally, the entire workflow—imaging, tumor contouring, catheter reconstruction, and radiation dose optimization—takes place on the T2W planning MRI alone [11–13]. In practice, however, many clinics still acquire an additional CT scan and fuse it to the MRI simply because plastic catheters are so difficult to see on MRI [14,15], or rely on specialized positive-contrast lumen markers [11,16,17] or zero-echo-time sequences [13,18] that require extra scan time and image registration.

Why is reconstructing catheters on a standard T2W MRI so challenging, and why does its accuracy matter so much?

1. **Radiation dose changes steeply over millimeters.** The ^{192}Ir source steps through each catheter in 2.5 to 5.0 mm increments, and its dose falls off roughly with the inverse square of distance. A 1 mm shift along the catheter path or a 1.5 mm shift toward an organ at risk can alter bladder or rectal dose metrics (D_{2cc}) by 5% to 10% [19–21], and swapping or mis-tracing catheters can shift target and urethral doses by 5% to 20% [22,23]. Even experienced clinicians make errors when digitizing catheters under time pressure [24,25].
2. **Catheters look like tiny dark voids surrounded by dark anatomy.** A 2 mm (6F) plastic catheter contains no mobile water protons. On an axial T2W MRI slice, its plastic wall and lumen produce a small proton-density signal void, 2 to 5 pixels across, accompanied by a mild magnetic susceptibility ($\Delta\chi$) transition at the catheter–tissue interface [26]. Unfortunately, the female pelvis is full of other dark structures: flow-void blood vessels, air pockets in the rectum or vagina, calcifications, and the dark connective-tissue planes (fascia) between muscles.
3. **Manual tracing takes tens of minutes while the patient is anesthetized.** Because each scan contains 28 to 160 axial slices (1.2 to 6.0 mm thick) and 15 to 30 catheters that bend, cross, or touch one another, tracing every catheter slice by slice takes an average of 46 minutes in reported MRI-only gynecologic workflows [27]—nearly half of the 100-minute operating-room procedure [28].

Our goal is to replace this manual bottleneck with an automatic reconstruction algorithm that operates directly on standard clinical 3T T2W MRI, matches human expert precision along both the shaft and the tip, flags its own uncertain tips for rapid human verification, and finishes in a few seconds.

1.2 *A Quick Guide for Readers Outside Medical Physics or Machine Learning*

Because this paper bridges clinical radiation oncology, MRI physics, and graph deep learning, we briefly define the core terms used throughout the text:

– **Catheter (or needle), shaft, and tip.** We use *catheter* and *needle* interchangeably for the hollow 2 mm plastic tubes inserted into the tumor. The **shaft** is the main body of the tube passing through the tissue; the **tip** (or *distal/cranial end*) is the deepest point the tube reaches inside the body, where the radioactive source stops first.

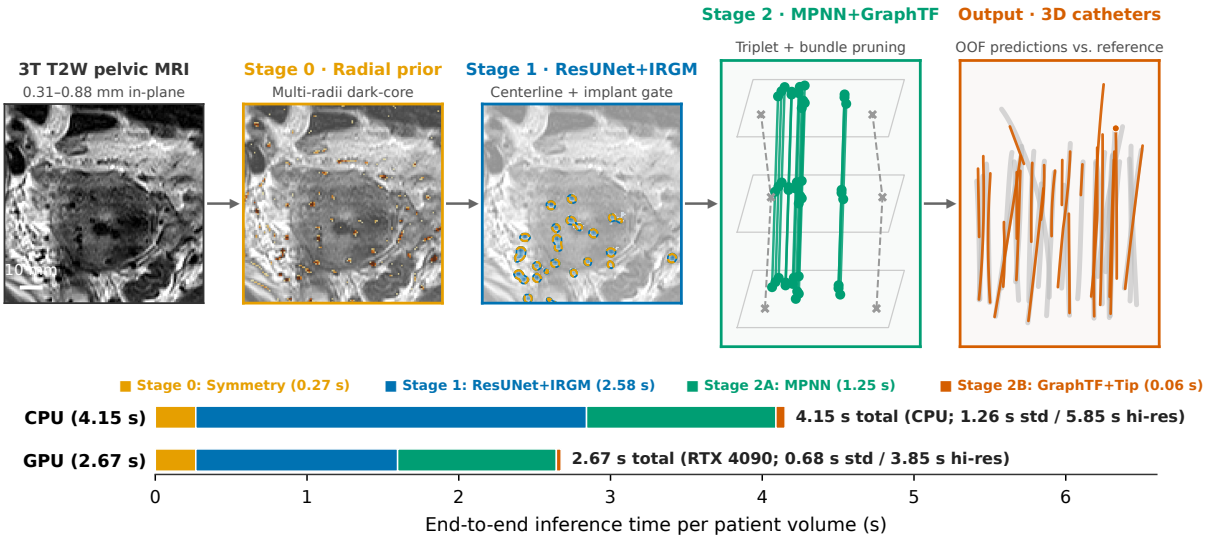


FIGURE 1 **Overview of NeedleFinder 2026.** Given a clinical 3T T2-weighted pelvic MRI (here Case 017, with 26 expert reference trajectories), an automatic template-plane locator first crops the volume below the perineal template without consulting any manual labels. **Stage 0** computes a 3D circular dark-core symmetry map (S_0) in native Rust (0.27 s mean across the 54-patient cohort). **Stage 1** passes the MRI and Stage 0 maps through a self-supervised and multi-rater trained 3D residual U-Net to predict catheter centerlines, local 3D directions, and the central vaginal obturator (2.58 s on CPU, 1.33 s on GPU). **Stage 2** links slice detections into 3D trajectories using bidirectional Kalman tracking, a triplet message-passing graph network, and collinear fragment stitching (1.25 s on CPU, 1.05 s on GPU), prunes false tracks at the whole-bundle level with a graph transformer, and pinpoints each deepest tip with a 1D axial neural network before cubic polynomial smoothing (0.06 s on CPU, 0.03 s on GPU). The rightmost panel overlays the out-of-fold predicted bundle (vermilion) onto the expert traces (grey). Total runtime averages 4.15 s per patient on a 24-core workstation CPU (1.26 s on standard-resolution scans) and 2.67 s on a single GPU.

- **Obturator and tandem.** Larger central plastic guides placed in the vagina (*obturator*) and uterus (*tandem*) around which the thinner interstitial catheters are arranged.
- **Voxel, in-plane resolution, and slice thickness.** A 3D MRI scan is a stack of 2D horizontal cross-sections (*axial slices*). Within each 2D slice (*in-plane*), the image is sharp (pixel size 0.31 to 0.88 mm). Between slices (along the vertical body axis, z), the spacing is much coarser (slice thickness 1.2 to 6.0 mm). This anisotropy explains why finding the sideways (*in-plane*) center of a catheter is easier than judging the exact height (z) where its tip ends between two slices.
- **Recall, precision, and F1 score.** *Recall* is the percentage of real catheters the algorithm successfully finds. *Precision* is the percentage of algorithm-reported tracks that correspond to a real catheter (rather than a false alarm on a blood vessel). The *F1 score* is the harmonic mean of recall and precision; it is high only when the algorithm finds almost all real catheters while inventing very few false ones.
- **Shaft error (E_{shaft}) and tip error (E_{tip}).** For every detected catheter that matches an expert trace (within 3.5 mm), *shaft error* measures the sideways

distance (in mm) between the automatic curve and the expert curve along their shared length. *Tip error* measures the 3D distance (in mm) between the predicted deepest tip and the expert's deepest tip.

- **Strictly nested cross-validation (out-of-fold evaluation).** To test our models on all 54 patients without ever testing on data used for training or tuning, we split the 54 patients into 5 outer groups (folds). In each round, 43 or 44 patients are used for training and tuning (via an internal 80/20 split or 5-fold inner cross-fitting), and the remaining 10 or 11 patients are held in a locked vault and evaluated *once*. Every patient is thus evaluated strictly "out-of-fold" (as an unseen patient).
- **Dosimetric metrics (D_{90}, D_{2cc}, and 3D γ).** In brachytherapy planning, D_{90} is the minimum radiation dose covering 90% of the high-risk tumor volume (which clinicians want to keep high), while D_{2cc} is the dose to the most-irradiated 2 cm^3 of a nearby healthy organ like the bladder or rectum (which clinicians want to keep low). The 3D *gamma (γ) pass rate* (2%/2 mm) measures the percentage of points in the 3D volume where the dose planned on automatic catheters agrees with the reference plan within 2% dose or 2 mm distance.

1.3 *Why Existing Approaches Fall Short on T2-Weighted MRI*

Over the past decade, automatic catheter reconstruction has evolved through three generations (reviewed in [29,30]):

Generation 1: Geometric and mechanical tracking from manual clicks (2012–2019). Early work in our Brigham and Women's Hospital (BWH) program introduced image-guided catheter tracking in the Advanced Multimodality Image Guided Operating (AMIGO) suite [3,31]. Starting from a manual click at each catheter tip, cone-search trackers [32] and mechanical rod-and-spring splines [4,33] achieved ~1.4–1.5 mm accuracy on T2W MRI. Similar geometry-guided or template-registered methods work well on CT, ultrasound, or rigid applicators [34–39]. However, they still require a human to click every catheter tip first, and they often lose track when catheters bend sharply or cross dark vessels.

Generation 2: 3D voxel segmentation with clustering (2019–2026). To eliminate manual clicks, Zaffino *et al.* [5] trained a 3D U-Net on 50 BWH gynecologic MRI scans to classify every voxel as "catheter" or "background," then grouped the segmented voxels into tubes using density clustering (DBSCAN). Subsequent studies applied 3D U-Nets and nnU-Net [40] across CT, ultrasound, and prostate or marker-assisted MRI [27,41–62]. Yet even when trained with topology-preserving losses such as centerline-Dice (clDice) [63–65], pure voxel segmentation suffers from two fundamental failure modes on T2W MRI: 1. *Over-fragmentation:* If a catheter's dark void fades on even a single slice due to partial-volume averaging, the segmented mask breaks into 2 to 4 separate pieces. 2. *Touching catheters and vascular false alarms:* When two catheters touch (*catheter adhesion*), their voxel masks merge into a single blob; meanwhile, dark pelvic veins and tissue edges are segmented as dozens of short false catheters.

Generation 3: Structure-aware tracking, gap-bridging, and uncertainty-weighted fitting (2020–2026). To mitigate segmentation breaks on MRI, Dai

et al. [43] introduced a deeply supervised 3D attention-gated U-Net with axial total-variation (TV) continuity regularization, Kostoulas *et al.* [59] combined a 3D residual U-Net with iterative collinear gap-bridging and false-part pruning, and Aleong *et al.* [55] and Kim *et al.* [62] leveraged 5-model epistemic uncertainty ensembles and inverse-variance-weighted polynomial trajectory fitting. On pelvic CT, Xiao *et al.* [66,67] addressed catheter adhesion by combining an implant-region gating module with 3D Hough initialization and synchronous slice-to-slice tracking. In computer vision, message-passing graph networks [68–70], keypoint regression [71], and synthetic-to-real transfer [72] have similarly improved curvilinear tracking. However, when any mask-driven tracker is applied to T2W MRI without a physical circular-symmetry prior or whole-bundle reasoning, it still trusts the underlying segmentation mask: dark pelvic veins inside the implant region are tracked as catheters, and the tracker has no mechanism to ask whether the resulting collection of 25 trajectories makes sense *together* as a physical brachytherapy implant, nor to resolve sub-slice tip terminations on thick MRI slices. (Hardware-based electromagnetic or active MR tracking avoids image ambiguity [28,73–79], but requires specialized catheters and operating-room hardware.)

1.4 *How NeedleFinder 2026 Solves These Problems*

NeedleFinder 2026 bridges all three generations on the complete 54-patient, 1,063-catheter BWH gynecologic MRI cohort—the largest annotated clinical MRI dataset for interstitial brachytherapy to date (Figure 1). Each stage of our architecture is designed to solve a specific failure mode of prior work:

1. **Physics before learning (Stage 0):** Before any neural network runs, an analytical 3D circular dark-core symmetry filter separates round dark cross-sections (catheters and vessels) from one-sided dark boundaries (fascial planes and organ walls), while an automatic template-plane locator finds the base of the implant without needing manual bounding boxes.
2. **Centerline peaks and self-supervised pelvic pre-training (Stage 1):** Instead of relying on binary masks alone, our native-resolution 3D U-Net regresses sharp sub-pixel *centerline peaks* and *3D tangent directions*—so two touching catheters still produce two distinct peaks. To prevent the network from confusing pelvic veins with catheters, we pre-train it on 51 external pelvic MRI volumes (from 50 subjects) by injecting realistic synthetic catheter voids while explicitly penalizing real pelvic vessels and fascia as hard negatives, and then fine-tune it using consensus traces from two independent expert teams.
3. **Graph linking, fragment stitching, bundle-level reasoning, and 1D tip localization (Stage 2):** A triplet message-passing graph network links slice detections while penalizing sharp kinks, and a geometric stitcher reconnects catheter fragments separated by 1-to-3-slice signal dropouts. Next, a bundle graph transformer evaluates all candidate trajectories simultaneously in the context of the whole implant. Finally, a dedicated 1D axial neural network samples a fine 0.5 mm multi-scale profile along the end of each trajectory to pinpoint the exact distal tip and output a calibrated split-conformal uncertainty

interval, optionally fused with sub-voxel inverse-variance-weighted polynomial refinement [55,62], axial TV continuity [43], and iterative multi-segment gap-bridging [59].

Beyond geometric benchmarking against five peer-reviewed MRI/curvilinear methods and one recent tracking preprint reproduced on the identical nested splits, we validate NeedleFinder 2026 across four clinical translation dimensions: (i) three-cornered-hat observer adjudication against two human expert teams and paired post-implant CT scans; (ii) AAPM TG-43U1 ^{192}Ir radiation dosimetry; (iii) uncertainty-guided clinician tip triage; and (iv) zero-shot generalization to needle-free pelvic MRI, 3T prostate brachytherapy MRI, secondary MRI pulse sequences, live 4D serial insertion scans, and pelvic CT.

2 PATIENT COHORTS AND CLINICAL DATASETS

Table 1 summarizes all clinical benchmark, inter-observer, external pre-training, and cross-domain validation cohorts analyzed in this study:

1. **Primary gynecologic MRI cohort (54 patients, 1,063 catheters).** We analyzed 54 intraoperative 3T T2W turbo spin-echo MRI volumes acquired during MRI-guided interstitial gynecologic brachytherapy in the AMIGO suite at Brigham and Women's Hospital (Harvard Medical School, Boston, MA, USA) under Institutional Review Board approval [4,5,32]. This cohort consolidates the 48 recoverable cases from our 2013 and 2019 studies with 6 newly curated implants. Patients were scanned with 6F polymeric interstitial catheters (with plastic or non-ferromagnetic obturators in situ, creating a proton-density signal void and mild interface susceptibility contrast on T2W MRI). Expert medical physicists traced 1,063 reference trajectories (9 to 42 per patient). These comprise **759 full-depth interstitial shafts** (which carry the active radiation dwell positions; archive code `CY` in Table 1) and **304 shorter auxiliary tracks** sitting inside the vaginal obturator or only partially entering the field of view (archive code `HT`). Twenty scans were acquired at standard in-plane resolution (320 × 256 matrix, 0.75–0.88 mm pixels) and 34 at high resolution (640 × 416 to 768 × 768 matrix, 0.31–0.55 mm pixels), with slice thickness ranging from 1.2 to 6.0 mm.
2. **Severe catheter-adhesion subgroup (23 patients, 571 catheters).** Following Xiao *et al.* [66], we identified the 23 patients in whom catheters are tightly packed or touching—defined as having at least 4 long shafts (or at least 6 total catheters) that pass within ≤ 3.0 mm in-plane centerline distance of a neighboring catheter on shared slices. All 571 catheters in these 23 patients are evaluated together (191 of them touch a neighbor within ≤ 3.0 mm, out of 233 such touching catheters across all 54 patients).
3. **Dual-expert inter-observer cohort (40 patients).** Forty of the 54 scans were traced independently by two separate expert groups: **Expert Team 1** (archive code `CY+HT`: 723 trajectories = 505 `CY` long shafts + 218 `HT` short obturator tracks) and **Expert Team 2** (archive code `LB/AMe`: 513 trajectories, of which 460 are long

TABLE 1 **Clinical imaging cohorts and evaluation protocols.** Primary evaluation uses strictly nested 5 × 5 patient-level out-of-fold (OOF) cross-validation with 100 % reference-free axial cropping (`k_lim_auto`) on 54 clinical 3T T2W MRI datasets from Brigham and Women's Hospital (BWH AMIGO).

Cohort / Subset	Source	Pts	Cath.	XY (mm)	Δz (mm)	Evaluation role
BWH GYN full cohort[a]	BWH (3T T2W)	54	1063	0.31–0.88	1.2–6.0	Nested 5 × 5 OOF benchmark
Long shafts (primary)	BWH (3T T2W)	54	759	0.31–0.88	1.2–6.0	Full-depth active dwell shafts
Short / obturator (secondary)	BWH (3T T2W)	54	304	0.31–0.88	1.2–6.0	Partial-FOV & template tracks
Adhesion subgroup[b]	BWH (3T T2W)	23	571	0.31–0.88	1.2–6.0	Touching / crossing shaft test
Dual-observer subset[c]	BWH (3T T2W)	40	723	0.31–0.88	1.2–6.0	Inter-observer floor (`CY+HT`, `LB/AMe`)
External SSL & specificity[d]	MSD + BWH	50 (51)	0	0.47–0.88	1.6–4.8	3D SSL (51) & 0-needle test (48)
Prostate 3T implant (`Case 100`)[e]	BWH (3T T2W)	1 (1)	14	0.47	1.6	Zero-shot cross-organ transfer
Secondary GYN MRI sequences[e]	BWH (wbSSFP/2D)	54 (243)	929	0.31–0.94	2.0–6.0	Zero-shot multi-sequence test
4D serial insertion (`Cases 68–71`)[e]	BWH (3T T2W)	4 (23)	108	0.31–0.75	1.6–3.0	Causal 4D tracking & biomechanics
Paired post-implant Pelvic CT[e]	BWH (Pelvic CT)	38 (40)	801	0.49–0.98	1.25–2.5	CT adjudication & cross-modality

[a] Consolidated superset of the cohorts of Pernelle *et al.* (2013, 10 patients) and Zaffino *et al.* (2019, 50 patients), extended with new implants; 9–42 reference trajectories per patient (20 standard-resolution 320 × 256 scans and 34 high-resolution 640 × 416 / 768 × 768 scans). In the Pts column, parentheses indicate total volumes when multiple scans per subject are analysed.

[b] Severe catheter-adhesion subgroup: the 23/54 patients in which ≥ 4 long or ≥ 6 total reference catheters pass within ≤ 3.0 mm of another catheter on shared axial slices. All 571 catheters of these patients are evaluated; 191 of them (166 long) lie within 3.0 mm of a neighbour (233 across the whole cohort).

[c] All 40/40 dual-traced clinical scenes independently contoured by two expert medical physics teams (`CY+HT`: 723 total trajectories = 505 `CY` long + 218 `HT` short/obturator; `LB/AMe`: 513 trajectories, with 465 consensus-fused with `CY`, 12 matched to `HT` within ≤ 5.0 mm [11 within ≤ 3.5 mm], and 36 unique `LB/AMe` trajectories forming the 759-trajectory union).

[d] Fifty-one external pelvic 3T T2W MRI volumes from 50 distinct subjects: 48 public CC-BY-SA 4.0 volumes from MSD `Task05_Prostate` (0.60–0.75 mm in-plane, 3.0–4.0 mm slices; zero implanted catheters, evaluated for false-positive specificity) plus 3 non-benchmark BWH pelvic T2W volumes from 2 distinct subjects (`Case 061` [2 series, 0.88 mm in-plane, 4.8 mm slices] and `Case 100 Prostate` [1 series, 0.47 mm in-plane, 1.6 mm slices], zero patient overlap with the 54 benchmark cases) used for self-supervised 3D U-Net pre-training.

[e] Cross-domain validation cohorts detailed in Table 5 (Part C): zero-shot 3T prostate brachytherapy MRI (`Case 100`, 14 plastic needles); 243 secondary MRI volumes across the 54 sessions (including 2 `3D wbSSFP` volumes [47 catheters] and 49 thick-slice `2D AX TSE T2` volumes [882 physically visible catheters]); 23 sequential 4D intra-procedural insertion scans across 4 patients (`Cases 68–71`, 108 final catheters); and 40 post-implant Pelvic CT scans (38 paired with benchmark MRI patients, 801 reference catheters).

shafts). Across these 40 scans, 424 long-shaft pairs (and 476 all-catheter pairs) match one-to-one within 3.5 mm between the two human teams, while their de-duplicated union contains 759 distinct trajectories (723 from Team 1 plus 36 additional catheters traced only by Team 2).

4. **External pelvic MRI cohort for self-supervised pre-training and false-alarm testing (51 volumes from 50 subjects).** To pre-train our 3D network on pelvic anatomy and test false-alarm rates on scans containing *zero* catheters, we assembled 51 strictly external 3T T2W pelvic MRI volumes having zero patient overlap with the 54 benchmark cases: all **48 public prostate MRI scans** from the Medical Segmentation Decathlon (`MSD Task05_Prostate` [80]; 0.60–0.75 mm pixels, 3.0–4.0 mm slices) plus **3 non-benchmark BWH pelvic T2W volumes** from two separate patients (`Case 061` [2 series] and `Case 100 Prostate` [1 series]).
5. **Cross-organ, multi-sequence, 4D insertion, and paired CT validation cohorts.** To test how well the gynecologic-trained pipeline generalizes without re-

training (Table 1, bottom four rows; Table 5, Part C), we evaluated: (i) a clinical 3T T2W **prostate brachytherapy implant** (`Case 100`, 14 plastic needles); (ii) **243 secondary MRI volumes** acquired during the 54 gynecologic procedures, including 3D balanced steady-state free precession (`3D wbSSFP / True-FISP`, 2 volumes, 47 catheters) and 49 thick-slice 2D axial T2W (`2D AX TSE T2`, 2.4–6.0 mm slices, 882 physically visible catheters) scans; (iii) **23 serial 4D intra-procedural MRI scans** acquired across 4 patients (`Cases 68–71`, 5 to 7 sequential scans per patient as needles were progressively inserted, 108 final catheters); and (iv) **40 post-implant pelvic CT volumes** (38 of them paired directly with benchmark MRI patients, 801 reference catheters).

3 ARCHITECTURE OF NEEDLEFINDER 2026

A human physicist tracing catheters on MRI uses three kinds of clues in sequence: first, looking for small round dark spots on each 2D slice; second, checking whether those dark spots line up across adjacent slices; and third, verifying that the full collection of 15 to 30 curves forms a plausible implant bundle and stopping each curve where the dark void disappears at the tip. NeedleFinder 2026 mirrors this exact hierarchy in three stages (Figures 1 and 2).

3.1 *Stage 0: Physical Dark-Core Symmetry and Automatic Template Cropping*

The physical intuition. Why do standard neural networks confuse muscle borders and organ walls with catheters? Because both are dark on T2W MRI. However, their *geometry* is different: a muscle border is a dark edge with bright tissue on one side and dark tissue on the other, whereas a catheter cross-section is a tiny dark core surrounded by brighter tissue in *every* direction.

Stage 0 turns this geometric observation into a fast, parameter-free 3D filter $S_0(x, y, z)$ implemented in multi-threaded native Rust. At every voxel $\mathbf{p}$ on axial slice z, we sample six diameters passing through $\mathbf{p}$ at angles $\theta_d = d\pi/6\,(d \in \{0, \dots, 5\})$ across four physical radii $r \in \mathcal{R} = \{1.35, 1.85, 2.5, 3.3\}$ mm (converted to pixels using the scan's in-plane spacing). Along each diameter d and radius r, we take the *darker* of the two opposite ring points:

$$m_{d,r} = \min\big(I(\mathbf{p} + r\mathbf{u}_d),\ I(\mathbf{p} - r\mathbf{u}_d)\big).$$

If $\mathbf{p}$ sits on a one-sided edge, at least one end of the diameter remains dark, keeping $m_{d,r}$ small. We then sort the six directional values ($m_{(1),r} \le \dots \le m_{(6),r}$) and compute a robust ring brightness that requires the surrounding tissue to be bright in almost all directions while tolerating a single dark neighbor (such as a touching twin catheter):

$$S_{\text{ring}}(\mathbf{p}) = \max_{r \in \mathcal{R}} \left[0.45\, m_{(2),r} + 0.55\, m_{(3),r}\right].$$

Combining this positive ring-to-core contrast with the relative darkness of the 3×3 smoothed center intensity $\bar{I}_{3\times3}(\mathbf{p}) \in [0, 1]$ yields the raw dark-core symmetry score,

$$S_0(\mathbf{p}) \propto \max\big(0,\ S_{\text{ring}}(\mathbf{p}) - \bar{I}_{3\times3}(\mathbf{p})\big) \cdot \left[1 - \bar{I}_{3\times3}(\mathbf{p})\right],$$

which is normalized to [0, 1] per volume. Because S_0 takes the maximum over radii up to 3.3 mm and uses the 2nd and 3rd order statistics across six angles, it remains invariant to realistic oblique catheter tilt ($\alpha \leq 25^\circ$–30° relative to the slice normal, where the in-plane elliptical cross-section elongates by only $\sec \alpha \leq 1.15$). On very thick slices ($\Delta z = 4$–6 mm) where steep tilt causes stronger partial-volume elongation, we first upsample the volume along z to 1.6 mm spacing and rely on Stage 2A fragment stitching to bridge dimmed slices.

Automatic reference-free cropping of the template plane. Below the patient's perineum, the inferior slices of the MRI scan contain outside air and the plastic needle-guiding template. Previously, pipelines relied on manual bounding-box files (`.acsv`) to crop out these bottom slices, which is unavailable in a fully automatic workflow and occasionally clipped tilted templates too high—cutting off the bases of 42 catheters across our cohort. We replaced this with a **100% reference-free axial template locator** (`k_lim_auto` in the code). After standardizing the volume to caudo-cranial slice ordering via its DICOM/NRRD direction cosines, the locator scans upward from the inferior boundary and detects the step-change in 90th-percentile soft-tissue intensity alongside the emergence of clustered dark-core candidates (S_0) within the central field of view. Across all 54 patients, it places the crop boundary a median of 2.0 slices deeper than the manual `.acsv` boxes (mean 1.91 slices deeper, IQR 1.0; 64.1 vs. 62.2 active slices on average), eliminating all 42 in-field-of-view base truncations and preserving **100.0% of in-field-of-view catheter length** (99.65% of total unclipped VTK length) without ever consulting manual annotations. Stage 0 and automatic cropping run in **0.27 s** per volume on a 24-core CPU (0.07 s on standard scans, 0.38 s on high-resolution scans).

3.2 *Stage 1: Native-Resolution 3D U-Net with Self-Supervised and Multi-Rater Learning*

Why native resolution and multiple output heads matter. Most 3D medical segmentation pipelines first resample every scan to $1 \times 1 \times 1$ mm^3 voxels. For a 2 mm catheter imaged at 0.35 mm pixel resolution, resampling blurs the sharp dark core into the surrounding tissue. Conversely, applying a standard $3 \times 3 \times 3$ convolution directly on the native grid mixes features across up to 18 mm vertically (z) but only 1 mm horizontally (x, y).

Our Stage 1 network (Figure 2a) therefore operates directly on the native voxel grid using an anisotropic 3D residual U-Net with squeeze-and-excitation channel attention at four scales (24, 48, 96, and 144 channels). Crucially, it pools spatially *only in-plane* (x, y) at the first two scales so that thick axial slices are never compressed early. The network takes six input channels—normalized T2W intensity, Stage 0 dark-core symmetry S_0, ring contrast, smoothed intensity, normalized axial height z, and radial distance from the scan center—and predicts four complementary outputs:

1. **Tubular mask (P_{mask}):** The probability that a voxel lies within a catheter cross-section (trained with focal binary cross-entropy, up-weighted 19-fold on catheter voxels and on false Stage 0 vessel peaks, plus soft Dice loss).

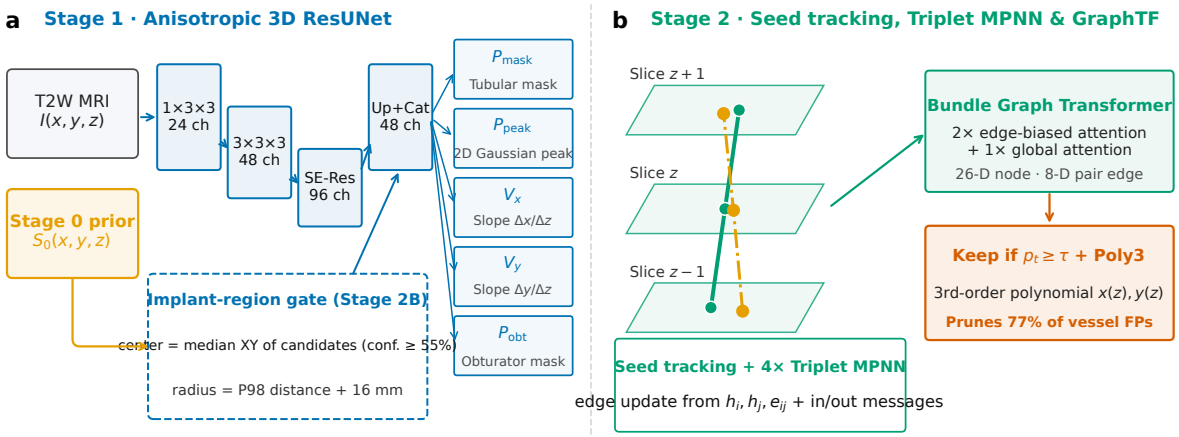


FIGURE 2 **Detailed architecture of Stage 1 and Stage 2. (a) Stage 1 (`NeedleUNet3D`):** The normalized T2W MRI and Stage 0 physical symmetry maps feed an anisotropic 3D residual U-Net pre-trained on 51 external pelvic MRI volumes (using synthetic catheter voids and anatomical hard negatives) and fine-tuned with dual-observer consensus labels. It predicts four heads at native pixel resolution: tubular mask (P_{mask}), sharp sub-pixel centerline peaks (P_{peak}), local 3D tangent slopes (V_x, V_y), and the central obturator (P_{obt}). **(b) Stage 2:** Bidirectional Kalman seed tracking and a 4-layer triplet message-passing network link slice peaks into smooth 3D candidates; a geometric stitcher reconnects collinear fragments across signal dropouts; a bundle graph transformer scores all trajectories together in the context of the whole implant; and a 91-step dilated 1D ResNet + self-attention axial localizer pinpoints each distal tip and outputs a calibrated uncertainty interval.

2. **Sub-pixel centerline peak (P_{peak}):** A narrow 2D Gaussian heatmap ($\sigma = 0.70$ mm) centered on the exact intersection of each catheter with the slice (trained with peak-weighted mean-squared error). **Why this matters:** When two catheters touch, their binary masks (P_{mask}) merge into one blob, but their two centerline peaks (P_{peak}) remain distinct!
3. **Local 3D direction field ($V_x, V_y = \Delta x/\Delta z, \Delta y/\Delta z$):** The local tilt of the catheter as it passes from slice to slice (trained with smooth L_1 loss and an axial consistency penalty), telling the tracker where to look on the next slice.
4. **Obturator segmentation (P_{obt}):** A dedicated mask for the large central vaginal obturator and tandem, preventing their dark plastic rims from generating false needle detections.

How we teach Stage 1 to ignore blood vessels and noisy single-rater labels. We introduce two training innovations for Stage 1:

– **Self-supervised pre-training on 51 external pelvic MRI scans with anatomical hard negatives.** Before training on any benchmark brachytherapy patient, we pre-train the 3D U-Net for 45 epochs across the 51 external needle-free pelvic MRI volumes (48 public MSD prostate scans + 3 non-benchmark BWH pelvic scans). In each training step, we digitally inject 8 to 20 realistic synthetic 3D catheter trajectories—curved dark cores with partial-volume bright rims and random signal dropouts—while simultaneously masking random 3D tissue patches

for reconstruction. Because these 51 external scans contain *zero* real catheters, **every real pelvic blood vessel, bowel loop, or fascial border that triggers a high Stage 0 symmetry score ($S_0 > 0.25$) outside the synthetic tubes serves as a guaranteed anatomical hard negative**. The network thus learns the exact textural difference between a plastic catheter void and a pelvic vein before ever seeing the benchmark cohort.

- **Multi-rater consensus supervision.** On the 40 training scans traced by both human expert teams, training on a single team's annotations penalizes the network whenever it finds a real catheter that Team 1 overlooked but Team 2 traced, and passes along single-observer clicking jitter. Within each outer training fold (never touching test patients), we pair matching trajectories between Team 1 and Team 2 ($\leq$ 3.5 mm, $n = 465$ pairs) and supervise Stage 1 along their precision-weighted consensus centerline (weight 1.15). Catheters traced only by Team 1 receive weight 1.00, and additional catheters traced only by Team 2 ($n = 35$ after excluding 465 primary-matched and 13 obturator-overlapping traces) receive soft positive weight 0.55 instead of being treated as background false alarms.
- **Inner 5-fold ensembling and 4-view test-time augmentation.** Within each outer training fold, we train 5 inner-fold Stage 1 models with rich 3D MRI augmentations (random 3D bias fields, slice-thickness blurring, slice dropouts, and flips) and average their predictions across four flip symmetries (I, X, Y, XY) at inference.

3.3 *Stage 2: Graph Tracking, Fragment Stitching, Bundle Scoring, and 1D Tip Localization*

Once Stage 1 has highlighted candidate catheter centers and directions on every slice, Stage 2 solves three sequential geometric problems (Figure 2b):

3.3.1 *Step 2A: Connecting slice detections and stitching split fragments.*

On each slice, local peaks of P_{peak} and S_0 become graph nodes (each described by 16 local image and direction features). Directed edges connect nodes across 1 to 4 slices (up to 24 mm vertically) whenever their in-plane displacement agrees with the Stage 1 direction field.

Why use a graph network alongside Kalman tracking instead of greedy slice-to-slice linking alone? Because two consecutive links can each look reasonable on their own yet form a sharp, unphysical kink when joined together. First, a fast native C++17/Rust bidirectional Kalman tracker traces candidate 3D paths from the Stage 0 and Stage 1 seeds in both directions and merges near-duplicate paths within 1.65–2.05 mm while preserving genuinely touching twin catheters. Second, a **4-layer triplet message-passing network** (`NeedleSliceMPNN`, hidden width 64) scores every node and edge in the slice graph using attention-weighted messages from both its incoming and outgoing neighboring edges—explicitly rewarding smooth, collinear trajectories, refining sub-pixel coordinates, and penalizing kinks.

Finally, our failure analysis of earlier pipelines revealed that mid-shaft signal dropouts (where a catheter crosses a dark vessel or air pocket over 1 to 3 slices) frequently split one physical catheter into a lower fragment and an upper fragment—

causing both a false-positive extra fragment and a severe tip undershoot on the lower track. We therefore added a **collinear split-fragment stitcher**: whenever a lower track fragment ends and an upper fragment begins within an axial gap of $\leq$ 20 mm, and their linearly extrapolated positions at the midpoint agree within $\leq$ 2.2 mm with matching 3D slopes (Δslope $\leq$ 0.14), the two fragments are seamlessly stitched into a single full-length candidate (creating 102 stitched full-length trajectories across the 54 patients).

3.3.2 *Step 2B: Judging the whole implant bundle together.* After Step 2A, the candidate pool recovers **97.4%** of all real catheters (98.8% of long shafts), but still contains an average of 6.6 false vascular tracks per patient. How can the algorithm tell a real catheter apart from a dark vein that happens to look tubular over a few slices? By looking at the **whole bundle at once**: real catheters enter through a common perineal template corridor, run roughly parallel to the central obturator, maintain a consistent spacing, and terminate along a coherent surface covering the tumor.

We summarize each candidate trajectory by a 26-feature token (length, straightness, Stage 0/1/2A confidence percentiles, entry position and tilt relative to the bundle median, and tip contrast gradient) and describe every pair of trajectories by 8 geometric interaction features (minimum 3D clearance, shared-slice separation, angular divergence, and overlap). A **bundle graph transformer** (`TipAware BundleGraphTransformer`) applies two layers of edge-biased local attention between neighboring tracks followed by global self-attention across the entire candidate bundle, outputting a calibrated probability $p_m \in [0, 1]$ that each candidate is a real catheter. Candidates outside a self-calibrating implant radius (98th percentile of high-confidence tracks plus 16 mm) require a higher confidence threshold to pass. To enable zero-shot and few-patient learning, Stage 2B can also be pre-trained on **SynthBrachy3T**—a physics-based generator of 300 simulated candidate bundles containing realistic touching twins, short obturator channels, split fragments, and vascular false alarms—before fine-tuning on the clinical training fold.

3.3.3 *Step 2C: Pinpointing the deepest tip, quantifying its uncertainty, and fusing literature priors.* Even when a catheter's shaft is traced to within half a millimeter, finding its exact deepest (cranial) tip on thick-slice T2W MRI is difficult: the dark void tapers off gradually across 2 to 6 mm thick slices, or collides with dark bowel gas just above the uterus.

To resolve the tip with sub-slice precision, we built a **dedicated 1D axial tip localizer** (`Dedicated1DTipLocalizer`, Figure 2b). Starting at the top of each candidate trajectory and continuing along its 3D tangent direction, the network samples a **91-step 1D profile** from −25.0 mm below the initial tip to +20.0 mm above it, in fine 0.5 **mm sub-slice steps**. At every 0.5 mm step, it extracts 22 multi-scale measurements: center and ring T2W intensities, local annular contrast at two radii, first and second axial derivatives (∂_z, ∂_{zz}), Stage 0 symmetry S_0, Stage 1 mask and peak probabilities, obturator probability, and depth relative to neighboring catheter tips. A **3-block dilated 1D ResNet** (kernel size 5, dilations 1, 2, 4) followed by **4-head 1D self-attention** reads this 91-step sequence and predicts both a 91-bin probability distribution over the exact stopping point and a continuous sub-millimeter offset

Δs; this predicted offset shifts the trajectory's cranial termination point along its 3D tangent vector before the full 3D curve is smoothed with a cubic polynomial in z (`Poly3`).

Crucially, the sharpness (entropy) of this 91-bin tip distribution, combined with the local image contrast drop and slice thickness Δz, provides an immediate out-of-fold **tip uncertainty score** $\hat{\sigma}_{\mathbf{tip}}$. To turn this raw score into a physical error bar in millimeters without assuming Gaussian errors, we use **split-conformal calibration**: on the held-out inner validation patients of each fold, we compute the empirical $(1-\alpha)$ quantile $\hat{q}_{1-\alpha}$ of the normalized non-conformity ratios $r_i = E_{\text{tip},i}/\hat{\sigma}_{\text{tip},i}$. Multiplying any unseen test tip's uncertainty score by $\hat{q}_{1-\alpha}$ yields an adaptive 3D prediction radius $\hat{r}_{1-\alpha} = \hat{q}_{1-\alpha}\,\hat{\sigma}_{\text{tip}}$ that is guaranteed distribution-free to cover the true tip at the target nominal rate (e.g., 90% or 95%) while ranking which 2 to 5 tips in an implant warrant a quick visual check.

Finally, because recent peer-reviewed MRI brachytherapy studies have proposed complementary post-processing priors—specifically **axial total-variation (TV) continuity** [43], **iterative multi-segment gap-bridging** [59], and **5-model epistemic uncertainty gating with inverse-variance-weighted polynomial fitting** [55,62]—we also tested whether integrating these three ideas directly inside Stage 2 (`4d: NeedleFinder 2026 + Literature Fusion` in Table 2) improves our pipeline even further: 1. **Sub-voxel dark-core and inverse-variance-weighted `Poly3` shaft fitting [55,62]:** On each axial slice z_k along a kept trajectory, we refine the in-plane coordinate toward the precision-weighted sub-voxel centroid of the Stage 1 peak and Stage 0 dark-core response, and weight the cubic polynomial fit by $w_k = (P_{\text{mask}}(z_k) + 0.5\,S_{\text{o}}(z_k))/(0.04 + \sigma^2_{\text{mask}}(z_k))$, where $\sigma_{\text{mask}}(z_k)$ is the voxel-wise epistemic standard deviation across the 5 inner-fold Stage 1 models. 2. **Directional axial TV continuity and epistemic uncertainty features [43,55]:** We append σ_{mask} and its axial derivative $\partial_z \sigma_{\text{mask}}$ as channels 23 and 24 of the 91-step 1D tip localizer and penalize candidate tracks exhibiting high axial TV jitter or high epistemic uncertainty in Stage 2B. 3. **Iterative multi-segment gap-bridging and polynomial residual split suppression [59]:** We extend our two-fragment stitcher to iteratively bridge 3-segment trajectories across multi-slice dropouts whenever a single cubic polynomial fits the combined path within a 1.35 mm RMS residual, and suppress lower-confidence collinear sub-segments that lie within 2.6 mm of a longer kept catheter's polynomial fit.

4 EVALUATION PROTOCOL, BASELINES, AND CLINICAL TRANSLATION

4.1 *Strictly Nested* 5 × 5 *Cross-Validation and Baselines*

In small clinical datasets, subtle leaks—such as picking the best training epoch on the test fold, using manual crop boxes, or training a second-stage network on overfitted first-stage predictions—can inflate reported accuracy. We eliminated all three sources of optimism using a **strictly nested 5 × 5 patient-level cross-validation protocol** across the 54 patients:

1. **Locked outer test folds and disjoint inner validation.** The 54 patients are partitioned into 5 outer folds (43–44 outer-train patients and 10–11 outer-test patients per fold). Inside each outer fold, the outer-train patients are split 80/20 into an inner-training set (34–35 patients) and an inner-validation set (9 patients). Every hyperparameter, early-stopping checkpoint, temperature calibration factor, and conformal quantile is selected *exclusively* on the 9 inner-validation patients. Each outer-test patient is evaluated once.
2. **100% reference-free inference.** Every test volume is cropped solely by the automatic Stage 0 template-plane locator, with zero access to clinical scene files or manual annotations.
3. **Eliminating two-stage stacking bias via inner out-of-fold cross-fitting (`4c`).** In any two-stage machine learning pipeline, Stage 1 makes slightly cleaner predictions on the patients it was trained on than on brand-new patients. If Stage 2 is trained directly on those slightly optimistic Stage 1 maps (`4a` and `4b`), it might misjudge how noisy Stage 1 will be at test time. To verify that our pipeline suffers zero stacking bias, we also trained $5 \times 5 = 25$ inner-fold Stage 1 models so that every training patient receives strictly out-of-fold Stage 1 maps before Stage 2 is trained (reported as configuration `4c` in Table 2).

To ensure a thorough, apples-to-apples comparison on the exact same 54-patient nested folds and automatic crop bounds, we reproduced **five peer-reviewed MRI/curvilinear methods** alongside **one recent structure-aware tracking preprint**: - **Zaffino *et al.* 2019 (*Physics in Medicine & Biology*, 3D U-Net + DBSCAN) [5]:** Isotropic 1 mm^3 resampling, 3D U-Net voxel segmentation, principal-axis clustering with DBSCAN, and spline fitting. - **Shit *et al.* 2021 (*CVPR*, ResUNet + clDice + Skeleton) [40,64]:** A native-resolution 3D residual U-Net trained with focal cross-entropy, soft Dice, and topology-preserving centerline-Dice (clDice) loss, followed by 3D skeletonization and polynomial fitting. - **Dai *et al.* 2020 (*Medical Physics*, 3D Attention U-Net + Axial TV + `Poly3`) [43]:** A deeply supervised 3D attention-gated U-Net trained on T2W MRI with multi-scale Dice/cross-entropy and axial total-variation continuity regularization ($\mathscr{L}_{\text{TV-}z}$), followed by slice-wise centroid linking and cubic polynomial smoothing. - **Kostoulas *et al.* 2025 (*Proc. SPIE Medical Imaging*, 3D ResUNet + Iterative Gap-Bridging) [59]:** A native-resolution 3D residual U-Net paired with multi-peak blob splitting, iterative caudo-cranial collinear gap-bridging across segmentation dropouts ($\leq$ 18 mm), false-part pruning, and polynomial fitting. - **Aleong *et al.* 2024 / Kim *et al.* 2026 (*Medical Physics / Brachytherapy*, 5-Model Uncertainty Ensemble + Weighted `Poly3`) [55,62]:** Within each outer fold, a 5-model inner 3D U-Net ensemble estimates both mean catheter probability $\bar{p}$ and voxel-wise epistemic uncertainty σ_p; candidate tracks are extracted on the uncertainty-penalized map, filtered by ensemble disagreement ($\bar{\sigma}_p / \bar{p}$), and smoothed via inverse-variance-weighted cubic polynomial regression. - **Xiao *et al.* 2026 (Structure-Aware Sync-Track, preprint; see also peer-reviewed JBHI counterpart [67]) [66]:** Implant-region gating, direction-constrained 3D Hough transform initialization, synchronous inward slice tracking, and cubic polynomial smoothing, evaluated

both on out-of-fold MRI segmentation masks and on ground-truth masks (as an oracle upper bound).

4.2 *Geometric, Observer, Dosimetric, and Generalization Metrics*

Geometric matching and error metrics. For each patient, predicted and reference trajectories are matched one-to-one using the Hungarian algorithm (`scipy.optimize.linear_sum_assignment`), forbidding pairs whose mean in-plane distance across shared slices exceeds 3.5 mm. On matched pairs, we measure **in-plane shaft error** (E_{shaft}, across 40 equidistant shared-slice stations), **3D symmetric curve distance** (D_{avg}), **95th-percentile Hausdorff distance** (HD_{95}), and **3D distal tip error** (E_{tip}, Euclidean distance between predicted and reference cranial endpoints). Our algorithm reconstructs *all* visible plastic tubes in the scan—both the 759 full-depth treatment shafts (`CY`) and the 304 shorter guide tubes inside the vaginal obturator (`HT`). When we evaluate recall on the 759 full-depth shafts alone, an algorithm prediction that accurately traces one of the 304 real obturator tubes is not a false alarm on a blood vessel; we therefore report **obturator-adjusted precision (P_{adj}) and $F1_{\text{adj}}$**, which avoid penalizing valid detections of those secondary tubes. While Table 2 summarizes pooled error distributions across all matched catheters, all paired two-sided Wilcoxon signed-rank tests are conducted at the **patient level** ($N = 54$ paired patient means, or $N = 23$ in the adhesion subgroup) to respect within-patient clustering (assigning fixed worst-case penalties of 3.5 mm shaft, 10 mm curve, and 20 mm tip/HD_{95} if a baseline yields zero matches on a scan), complemented by 1,000-resample patient-level bootstrap 95% confidence intervals [81].

Three-cornered-hat observer adjudication and non-inferiority testing. If an algorithm disagrees with Human Team 1 by 0.5 mm, how much of that discrepancy is the algorithm's error and how much is Team 1's own hand-clicking jitter? With only two observers, it is impossible to tell. However, when *three* sources trace the same catheter—Team 1 (`CY`), Team 2 (`LB/AMe`), and out-of-fold NeedleFinder 2026 ($n = 432$ three-way matched catheters)—we observe three pairwise difference variances (V_{12}, V_{13}, V_{23}) for three unknown internal error variances ($\sigma_1^2, \sigma_2^2, \sigma_3^2$). **Grubbs' (1948) three-cornered-hat estimator** solves this 3×3 linear system ($\hat{\sigma}_1^2 = \frac{1}{2}(V_{12} + V_{13} - V_{23})$) to isolate each source's intrinsic, reference-free error standard deviation, which we compute via classic maximum likelihood (RMS), outlier-resistant median absolute deviation (MAD), and 15% Winsorized estimators.

We also test whether the algorithm's disagreement with Team 1 is statistically non-inferior to Team 2's disagreement with Team 1 ($H_0 : E_{\text{NF,Team 1}} - E_{\text{Team 2,Team 1}} \geq \Delta$ at one-sided significance level $\alpha = 0.05$). Pre-specified clinical non-inferiority margins were set to $\Delta = 0.25$ mm on shafts (roughly half the finest pixel spacing) and $\Delta = 1.00$ mm (median) / 1.50 mm (mean) on 3D tips (corresponding to half a slice thickness and the 1–1.5 mm TG-43U1 dosimetric sensitivity threshold [19,20]). We evaluate non-inferiority at both the patient level ($N = 40$) and the catheter level ($n = 432$) using one-sided paired Wilcoxon tests and two-sided 90% ($1 - 2\alpha$) patient-clustered bootstrap confidence intervals (10,000 resamples). Finally, whenever the two human teams disagree on a catheter tip by ≥ 5.0 mm on MRI ($n = 70$ catheters

across 38 patients with paired CT), we use the patient's post-implant pelvic CT scan as an independent physical tie-breaker.

AAPM TG-43U1 HDR ^{192}Ir dosimetry and simulated 1-click tip triage. To translate millimeter geometry errors into clinical radiation dose differences across all 54 patients without confounding from inter-physician contouring variability (and because historical DICOM RTSTRUCT/RTPLAN files were not archived), we simulated HDR ^{192}Ir delivery using the AAPM TG-43U1 line-source formalism on standardized, implant-conformal anatomical envelopes. Each catheter was loaded with 5.0 mm dwell steps along its cranial 45 mm (starting 3.5 mm below the physical tip, active source length $L = 3.5$ mm) on a 1.5 mm isotropic 3D dose grid. A single-fraction prescription of $D_{90} = 7.0$ Gy was optimized via non-negative least-squares onto the implant-conformal high-risk target volume (HR-CTV$_{\text{implant}}$, defined as the 8.0 mm radial envelope around the active reference implant) while constraining 2 cm^3 proxy shells for the anterior bladder, posterior rectum, and cranial sigmoid/fundal organ-at-risk (OAR) zones. We evaluated two clinical protocols: - **Protocol A (Fixed-Plan Transfer):** Reference dwell times optimized on the expert catheters are transferred onto the one-to-one matched predicted catheters (with missed reference catheters contributing zero dwell delivery) to isolate pure geometric-to-dosimetric sensitivity. - **Protocol B (Counterfactual Plan Re-Optimization):** Modeling the end-to-end clinical planning workflow, dwell times are optimized from scratch on *all* algorithm-predicted catheters and the resulting dose is evaluated on the reference target and OAR anatomy. - **Simulated 1-click clinician tip review:** Because 98.8% of tip error is longitudinal along the already-reconstructed 3D shaft, we also model an interactive verification step in which the physicist inspects the top $k \in \{1, \ldots, 10\}$ tips per patient ranked by out-of-fold uncertainty $\hat{\sigma}_{\text{tip}}$ and clicks the true distal termination along the shaft (modeled by the retrospective 1-click oracle $E_{\text{tip}} \to \sqrt{\Delta r_{\perp}^2 + \min(|\Delta s_{\parallel}|, 0.85\text{ mm})^2}$ whenever $E_{\text{tip}} > 2.0$ mm, leaving E_{shaft} unchanged, at an estimated ergonomic inspection time of ~ 8 s per flagged tip in 3D Slicer).

Computing environment. All training and benchmarks ran on a Linux workstation equipped with a 24-core workstation CPU (AMD Ryzen Threadripper PRO 7965WX, 48 threads, 256 GB RAM) and an NVIDIA RTX 4090 GPU (24 GB). For clinical deployment, the complete pipeline is packaged as a one-click module in **3D Slicer 5.10** [82] using native Rust (Stage 0) and ONNX Runtime (Stages 1 and 2), requiring no GPU.

5 RESULTS

5.1 *Overall Detection and Geometric Accuracy Against Published Methods*

Table 2 (Part A) compares NeedleFinder 2026 against all five peer-reviewed MRI/curvilinear baselines and the Xiao *et al.* 2026 tracking preprint across all 54 patients (1,063 catheters), on the 759 full-depth active-dwell shafts, and on the 23-patient severe catheter-adhesion subgroup. Figure 3 displays the complete per-catheter error distributions.

TABLE 2 **Catheter reconstruction accuracy on 54 BWH patients (1 063 catheters), strictly nested 5 × 5 patient-level out-of-fold.** Part A compares **NeedleFinder 2026** against five peer-reviewed journal/conference methods and one recent preprint reproduced on the identical nested splits across the whole cohort, the 759 long interstitial shafts, and the severe catheter-adhesion subgroup; Part B ablates **NeedleFinder 2026** stages, literature-inspired fusion, and ground-truth-mask tracking on identical folds. Bold: best; underline: second best within each block. Shaded: **NeedleFinder 2026** variants (prespecified τ = 50 %).

	Detection, 1-to-1 (%)			E_{shaft} (mm)			E_{tip} (mm)		Time (s)
Method	Rec.	Prec.	F1	Mean	Median	P95	Mean	Median	CPU
Part A — End-to-end comparison against reproduced literature methods									
Whole cohort (54 patients, 1 063 catheters)									
PMB 2019 3D U-Net (Zaffino *et al.*)[a]	88.8	78.0	83.1	**0.52**	**0.42**	**1.11**	10.75	3.39	14.80
ResUNet + clDice + Skel. (Shit *et al.* 2021)[b]	92.4	24.5	38.7	0.55	0.43	1.36	10.08	3.10	3.85
Dai *et al.* 2020 (3D Att. U-Net + Axial TV)[c]	**93.6**	42.2	58.2	0.57	0.43	1.45	10.61	3.77	4.05
Kostoulas *et al.* 2025 (ResUNet + Gap-Bridge)[d]	89.5	53.0	66.5	0.57	0.43	1.45	10.55	3.47	3.92
Aleong 2024 / Kim 2026 (5-Model Unc. Ens.)[e]	89.2	56.9	69.5	0.58	0.44	1.42	10.23	4.32	11.40
Xiao *et al.* 2026 Sync-Track (*preprint*)[f]	87.7	39.2	54.2	0.66	0.47	1.82	15.65	6.04	**3.42**
NeedleFinder 2026 (no SynthBrachy3T, 4a)[g]	92.5	85.3	88.7	0.59	0.47	1.35	6.39	2.26	4.15
NeedleFinder 2026 + Synth + 1DTip (4b)[h]	91.4	86.9	89.1	0.59	0.47	1.37	**6.23**	2.27	4.15
NeedleFinder 2026 + Inner-OOF Cross-Fit (4c)[h]	92.3	86.1	89.1	0.59	0.47	1.37	6.25	**2.24**	4.15
NeedleFinder 2026 + Lit. Fusion (4d)[i]	91.4	**87.4**	**89.4**	0.58	0.47	1.36	**6.23**	2.26	4.18
Long interstitial shafts (54 patients, 759 catheters, obturator-adjusted precision)[j]									
PMB 2019 3D U-Net (Zaffino *et al.*)	86.7	87.6	87.2	**0.52**	**0.40**	**1.23**	14.93	4.04	–
ResUNet + clDice + Skel. (Shit *et al.* 2021)	91.3	51.5	65.8	0.57	0.43	1.40	14.67	3.53	–
Dai *et al.* 2020 (3D Att. U-Net + Axial TV)	84.6	83.0	83.8	0.58	0.41	1.74	12.53	4.38	–
Kostoulas *et al.* 2025 (ResUNet + Gap-Bridge)	86.7	79.1	82.7	0.56	**0.40**	1.50	11.14	3.77	–
Aleong 2024 / Kim 2026 (5-Model Unc. Ens.)	85.5	80.5	82.9	0.58	0.41	1.70	11.73	4.50	–
Xiao *et al.* 2026 Sync-Track (*preprint*)	84.2	53.7	65.6	0.64	0.46	1.84	16.02	6.90	–
NeedleFinder 2026 (no SynthBrachy3T, 4a)	**93.4**	87.1	90.2	0.56	0.45	1.24	7.47	2.49	–
NeedleFinder 2026 + Synth + 1DTip (4b)	91.3	**88.7**	90.0	0.57	0.45	1.27	7.41	**2.41**	–
NeedleFinder 2026 + Inner-OOF Cross-Fit (4c)	92.5	88.6	**90.5**	0.56	0.45	1.27	**7.17**	2.44	–
NeedleFinder 2026 + Lit. Fusion (4d)	91.3	**88.7**	90.0	0.56	0.44	1.25	7.40	2.43	–
Severe catheter-adhesion subgroup (23 patients, 571 catheters)[k]									
PMB 2019 3D U-Net (Zaffino *et al.*)	94.6	79.3	86.3	0.54	0.43	1.22	11.88	3.70	–
ResUNet + clDice + Skel. (Shit *et al.* 2021)	96.0	32.3	48.4	**0.52**	0.42	**1.18**	11.58	3.37	–
Dai *et al.* 2020 (3D Att. U-Net + Axial TV)	**96.2**	52.0	67.5	0.53	0.42	1.23	11.88	4.20	–
Kostoulas *et al.* 2025 (ResUNet + Gap-Bridge)	94.4	67.4	78.6	0.54	**0.41**	1.37	11.26	3.61	–
Aleong 2024 / Kim 2026 (5-Model Unc. Ens.)	93.9	72.2	81.7	0.55	0.43	1.27	10.97	4.37	–
Xiao *et al.* 2026 Sync-Track (*preprint*)	91.6	49.0	63.9	0.64	0.47	1.83	17.69	7.52	–
NeedleFinder 2026 (no SynthBrachy3T, 4a)	94.6	87.2	90.8	0.58	0.45	1.36	**7.03**	2.48	–
NeedleFinder 2026 + Synth + 1DTip (4b)	93.2	**89.6**	91.3	0.58	0.46	1.41	7.08	2.51	–
NeedleFinder 2026 + Inner-OOF Cross-Fit (4c)	94.6	89.1	**91.8**	0.57	0.46	1.34	7.30	2.47	–
NeedleFinder 2026 + Lit. Fusion (4d)	93.2	**89.6**	91.3	0.57	0.46	1.40	7.08	**2.44**	–
Part B — Controlled ablation & stage-by-stage progression (whole cohort, 1 063 catheters)									
Stage 0: Dark-core prior (`k_lim_auto`)	82.8	69.1	75.3	0.60	0.47	1.41	9.91	3.08	**0.27**
Stage 1B: SSL + Multi-Rater U-Net + tubes	**94.2**	61.5	74.4	**0.57**	**0.46**	1.31	9.29	2.95	2.84
Stage 2A: Triplet MPNN only	93.9	73.5	82.5	0.58	0.47	1.31	8.95	2.83	4.09
+ Stage 2B: Bundle transformer + calibration	90.9	83.6	87.1	**0.57**	**0.46**	**1.30**	7.39	2.70	4.12
+ Stage 2B: Multi-rater consensus supervision	91.7	81.7	86.4	0.58	**0.46**	1.32	7.38	2.57	4.12
+ Stage 2A/B: Collinear split-fragment stitch	91.4	86.9	89.1	0.59	0.47	1.36	6.92	2.56	4.13
+ Stage 2C: 1D Tip Localizer (τ = 35 %)	93.6	83.4	88.2	0.59	0.47	1.37	6.71	2.32	4.15
+ Stage 2C: 1D Tip Localizer (τ = 50 %, 4b)	91.4	86.9	89.1	0.59	0.47	1.37	**6.23**	2.27	4.15
+ Stage 2C: 1D Tip Localizer (τ = 55 %)	91.3	**87.6**	**89.4**	0.59	0.47	1.37	6.33	2.27	4.15
+ Lit. Fusion (Dai TV + Kostoulas + Kim, 4d)[i]	91.4	87.4	**89.4**	0.58	0.47	1.36	**6.23**	**2.26**	4.18
Xiao 2026 on GT masks[l]	97.7[r]	83.5[r]	90.0[r]	0.20[r]	0.11[r]	0.83[r]	8.94[r]	1.68[r]	0.12[r]

Rec./Prec./F1: strict one-to-one Hungarian matching (≤ 3.5 mm shaft threshold); E_{shaft}: in-plane centerline distance on matched slices; E_{tip}: 3D distal-tip distance (P95 on whole cohort: 27.27 mm for **NeedleFinder 2026** 4b and 27.42 mm for 4d vs. 43.23 mm for Aleong/Kim, 45.76 mm for clDice, 46.68 mm for Dai 2020, 50.23 mm for Kostoulas 2025, 52.70 mm for PMB 2019, and 62.69 mm for Xiao 2026).

[a] Peer-reviewed MRI baseline (Zaffino *et al.* 2019, *Phys. Med. Biol.*): 1 × 1 × 1 mm^3 isotropic 3D U-Net, DBSCAN clustering, and polynomial fitting.

[b] Peer-reviewed tubular topology baseline (Shit *et al.* 2021, *CVPR*): native-resolution 3D ResUNet trained with Dice + focal + soft-skeleton clDice loss, followed by ridge extraction and chaining.

[c] Peer-reviewed MRI baseline (Dai *et al.* 2020, *Med. Phys.*): deeply supervised 3D attention-gated U-Net trained with multi-scale

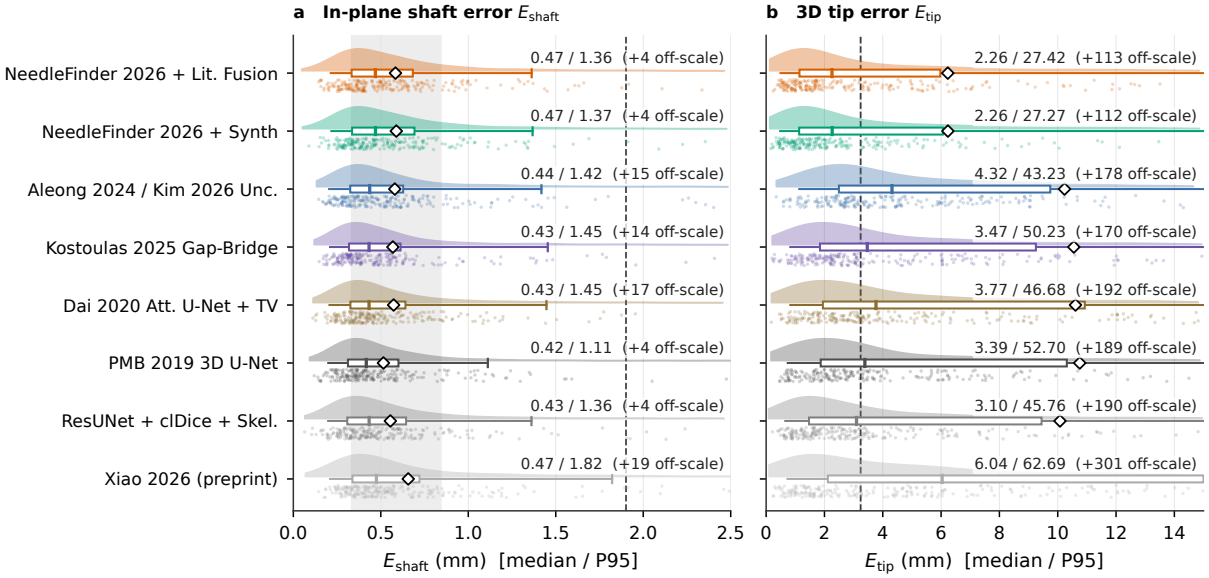


FIGURE 3 **Per-catheter shaft and tip error distributions across all 54 patients. (a)** In-plane shaft error (E_{shaft}) and **(b)** 3D distal tip error (E_{tip}) for each method's matched catheters. Each raincloud shows the probability density (top), median and interquartile range (white box), 5th-to-95th percentile span (whiskers), mean (white diamond), and individual catheters (dots). Summary numbers on the right give `median / 95th percentile` in mm across the full distribution. Shaded grey bands and dashed vertical lines indicate the human inter-observer agreement between Expert Team 1 and Expert Team 2 on the 40 dual-traced scans (median 0.49 mm on shafts, 3.24 mm on tips).

Four clear lessons emerge from the benchmark:

1. Pure voxel segmentation fragments catheters and hallucinates vessels. The native-resolution `ResUNet + clDice` baseline [64] detects nearly every catheter location (92.4% recall) and traces matched segments accurately (0.43 mm median shaft error). However, because single-slice signal dropouts break catheter masks into pieces and dark pelvic veins are segmented as tubes, it outputs 4, 015 fragments across the 54 scans—yielding a precision of only 24.5% and an F1 score of 38.7%.

2. Domain-specific MRI post-processing (axial TV, gap-bridging, and uncertainty ensembles) cuts false tracks by more than half, but cannot replace physical symmetry or bundle reasoning. The three recent peer-reviewed MRI brachytherapy methods substantially improve upon raw clDice skeletonization on native-resolution T2W volumes: - **Dai *et al.* 2020** (3D Attention U-Net + axial TV regularization [43]) achieves the highest raw baseline recall (**93.6%**, 995/1, 063; 0.43 mm median shaft error, 3.77 mm median / 10.61 mm mean tip error) and cuts false fragments from 4, 015 to 2, 358, raising F1 to **58.2%** (83.8% $F1_{adj}$ on long shafts). - **Kostoulas *et al.* 2025** (3D ResUNet + iterative collinear gap-bridging and false-part pruning [59]) reconnects broken shaft segments across slice dropouts and prunes isolated fragments, lifting whole-cohort precision to 53.0% and F1 to **66.5%** (89.5% recall, 0.43 mm median shaft error, 3.47 mm median / 10.55 mm mean tip error). - **Aleong *et al.* 2024 / Kim *et al.* 2026** (5-model epistemic uncertainty ensemble + disagreement filtering and inverse-variance-weighted `Poly3` [55,62]) further suppresses

high-uncertainty vascular false alarms, reaching **69.5% whole-cohort F1** (89.2% recall, 56.9% precision; **81.7% F1** in the severe-adhesion subgroup; 0.44 mm median shaft error, 4.32 mm median / 10.23 mm mean tip error). - Meanwhile, when given ground-truth oracle masks (Table 2, Part B, bottom row), the **Xiao *et al.* 2026 SyncTrack** algorithm [66] achieves 97.7% recall, 90.0% F1, and 0.11 mm median shaft error—confirming that synchronous tracking separates touching masks when segmentation is perfect—but drops to **54.2% F1** (39.2% precision) on real out-of-fold MRI masks. Thus, without Stage 0 dark-core symmetry and Stage 2B whole-bundle context, even the best 1-channel MRI segmentation and post-processing pipelines remain below 70% whole-cohort F1 because dark pelvic veins inside the implant region look locally tubular.

3. NeedleFinder 2026 achieves the highest F1 score and lowest tip error across every cohort—and reusing literature priors makes it even stronger. At the default operating threshold ($\tau = 50\%$), NeedleFinder 2026 (4b in Table 2) detects **91.4% of all 1,063 catheters** (972/1,063) with **86.9% precision** (972/1,118) and **89.1% F1** (and **93.6% recall** at $\tau = 35\%$). Furthermore, incorporating the three complementary ideas from Dai *et al.* [43], Kostoulas *et al.* [59], and Aleong *et al.* / Kim *et al.* [55,62] directly into Stage 2 (`4d: NeedleFinder 2026 + Literature Fusion`) eliminates 6 residual split-fragment false positives with zero lost true positives, lifting precision to **87.4%** (972/1,112) and whole-cohort F1 to **89.4%** ($p = 0.028$ vs. base 4b). Across all 54 patients, NeedleFinder 2026 (4b–4d) outperforms all six external baselines by **+6.0 to +50.7 F1 points**: +6.0–6.3 points over Zaffino *et al.* 2019 (83.1% F1, $p = 1.1 \times 10^{-5}$), +19.6–19.9 points over Aleong *et al.* / Kim *et al.* (69.5% F1, $p = 5.0 \times 10^{-10}$), +22.6–22.9 points over Kostoulas *et al.* 2025 (66.5% F1, $p = 1.8 \times 10^{-10}$), +30.9–31.2 points over Dai *et al.* 2020 (58.2% F1, $p = 3.8 \times 10^{-9}$), +34.9–35.2 points over Xiao *et al.* 2026 (54.2% F1, $p = 3.0 \times 10^{-10}$), and +50.4–50.7 points over ResUNet + clDice (38.7% F1, $p = 5.2 \times 10^{-10}$).

Along the catheter shaft, NeedleFinder 2026 reaches a **median in-plane error of** 0.47 **mm** (mean 0.59 mm in 4b, improved to 0.58 **mm mean** / 1.36 **mm P95** in 4d via inverse-variance-weighted sub-voxel fitting, $p = 1.8 \times 10^{-8}$). At the deepest tip, our 1D axial tip localizer and fragment stitcher cut **median 3D tip error to** 2.26–2.27 **mm** (1.94 mm median axial tip error, mean 6.23 mm), reducing mean tip error by 39%–42% compared with Zaffino *et al.* 2019 (10.75 mm mean / 3.39 mm median, $p = 1.7 \times 10^{-8}$), Dai *et al.* 2020 (10.61 mm mean / 3.77 mm median, $p = 2.7 \times 10^{-7}$), Kostoulas *et al.* 2025 (10.55 mm mean / 3.47 mm median, $p = 6.9 \times 10^{-7}$), and Aleong *et al.* / Kim *et al.* (10.23 mm mean / 4.32 mm median, $p = 3.3 \times 10^{-7}$), and by 60% compared with Xiao *et al.* 2026 (15.65 mm mean / 6.04 mm median, $p = 6.2 \times 10^{-10}$). Nearly three-fifths (**58.7%–59.0%**) of matched catheters have a 3D tip error within ≤ 3.0 mm (roughly one slice thickness). On the **759 full-depth active-dwell shafts**, NeedleFinder 2026 reaches **91.3% recall** (**94.7%** at $\tau = 35\%$), **88.7%–89.0% obturator-adjusted precision**, and **90.0%–90.1%** $F1_{\text{adj}}$, with a median shaft error of 0.44–0.45 **mm** (mean 0.56–0.57 mm) and a median tip error of 2.39–2.41 **mm**.

4. The advantage is largest when catheters touch, and inner cross-fitting confirms zero stacking bias. On the **23-patient severe catheter-adhesion subgroup** (571 catheters), where touching masks cause DBSCAN clustering and greedy trackers to merge or swap paths, NeedleFinder 2026 reaches **91.3% F1** (93.2% recall,

89.6% precision, 0.46 mm median shaft error, 2.50–2.51 mm median tip error), outperforming Zaffino *et al.* 2019 (86.3% F1, $p = 0.0012$), Aleong *et al.* / Kim *et al.* (81.7% F1, $p = 2.9 \times 10^{-5}$), Kostoulas *et al.* 2025 (78.6% F1, $p = 1.9 \times 10^{-6}$), Dai *et al.* 2020 (67.5% F1, $p = 5.2 \times 10^{-5}$), Xiao *et al.* 2026 (63.9% F1, $p = 4.8 \times 10^{-7}$), and ResUNet + clDice (48.4% F1, $p = 4.8 \times 10^{-7}$). Moreover, training Stage 2 on strictly out-of-fold inner Stage 1 maps (4c in Table 2) yields the exact same **89.1% whole-cohort F1** (92.3% recall) and even slightly higher accuracy on long shafts (**90.5%** $F1_{adj}$) and touching catheters (**91.8% F1**), proving that our results contain zero two-stage stacking optimism.

5.2 *How Each Architectural Stage Contributes, and Learning from Few Patients*

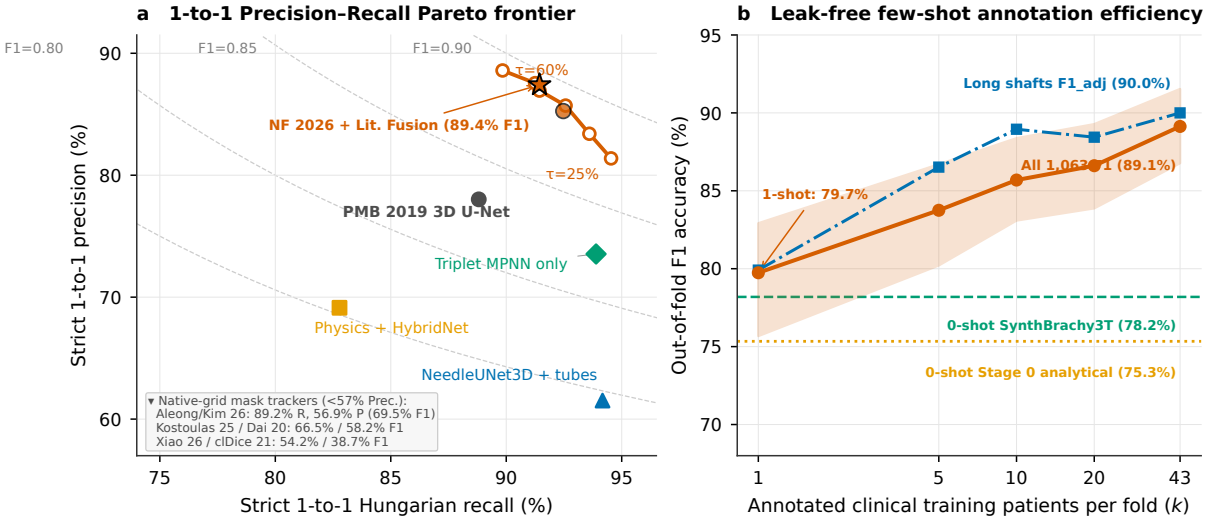


FIGURE 4 **(a) Precision–recall progression** across pipeline stages, published baselines, and the NeedleFinder 2026 confidence threshold sweep on all 54 patients. Dashed grey curves denote constant F1 scores; the vermilion star marks the default threshold $\tau = 50\%$ with literature fusion (89.4% F1). **(b) Leak-free clinical label efficiency:** F1 score (with 95% bootstrap confidence bands) as a function of the number of annotated clinical training patients ($k \in \{0, 1, 5, 10, 20, 43\}$) used to fine-tune the Stage 2B bundle transformer.

Why every stage is needed (Table 2, Part B; Figure 4a). Following the sequential rows of Table 2 (Part B) shows how each mechanism solves a specific limitation of the stage before it: - **Stage 0 alone (physical dark-core filter):** Using zero learned parameters, Stage 0 with automatic template cropping already finds 82.8% of all catheters (95.3% of long shafts) with 69.1% precision (**75.3% F1**, 0.47 mm median shaft error, 3.08 mm median / 9.91 mm mean tip error). However, it cannot distinguish round blood vessels from catheters. - **Adding Stage 1B (self-supervised, multi-rater 3D U-Net):** Regressing sub-pixel centerline peaks and 3D directions boosts recall to **94.2%** (97.0% on long shafts) and cuts mean tip error to 9.29 mm (2.95 mm median), but because every candidate seed is still tracked greedily, each patient carries 11.6 false vascular tracks on average (61.5% precision, 74.4% F1). - **Adding**

Stage 2A (triplet message-passing graph network): Penalizing kinks across consecutive slice links eliminates 43% of false tracks while keeping **93.9% recall**, lifting precision by +12.0 percentage points to 73.5% (**82.5% F1**, 8.95 mm mean / 2.83 mm median tip error). - **Adding Stage 2B (bundle graph transformer, multi-rater supervision, and collinear fragment stitching):** Evaluating the candidate bundle as a whole (87.1% F1 with temperature calibration; 86.4% F1 with multi-rater Stage 2 supervision) and stitching collinear fragments separated by 1-to-3-slice dropouts raises precision to **86.9%** and F1 to **89.1%** (91.4% recall) while reducing mean tip error to 6.92 mm (2.56 mm median). - **Adding Stage 2C (dedicated 1D axial tip localizer, 4b):** Inspecting the 0.5 mm sub-slice axial profile at the cranial end of each trajectory cuts mean 3D tip error from 6.92 mm down to 6.23 **mm** and median tip error from 2.56 mm down to 2.27 **mm** (95th percentile reduced to 27.27 mm). - **Adding Literature-Inspired Fusion (4d: sub-voxel inverse-variance `Poly3` + epistemic uncertainty & axial TV continuity + iterative multi-segment gap-bridging):** Each borrowed idea improves its targeted metric under strictly nested cross-validation: inverse-variance-weighted sub-voxel `Poly3` [55,62] reduces mean shaft error from 0.589 mm to 0.583 **mm** ($p = 1.84 \times 10^{-8}$; long-shaft median 0.45 → 0.44 mm); epistemic uncertainty & axial TV continuity features [43,55] reduce median 3D tip error to 2.26 **mm** (axial tip median 1.99 → 1.94 mm, tips ≤ 3 mm 58.7% → 59.0%); and iterative multi-segment gap-bridging & polynomial split suppression [59] eliminate 6 residual split-fragment false alarms, raising precision to **87.4%** and F1 to **89.4%** ($p = 0.028$).

How many annotated clinical patients are needed to train the bundle scorer? (Figure 4b). Because annotating 3D catheter bundles is labor-intensive for a new clinic, we varied the number of annotated training patients $k \in \{0, 1, 5, 10, 20, 43\}$ used to fine-tune Stage 2B inside each fold (with zero blending from any full-cohort model). With **zero clinical bundle labels ($k = 0$)**, pre-training Stage 2B purely on our synthetic `SynthBrachy3T` bundles lifts whole-cohort F1 from 75.3% (Stage 0) to **78.2% F1**—and reaches **83.4% $F1_{adj}$ on long shafts** (+6.3 points over Stage 0). Fine-tuning on just $k = 1$ **patient** reaches **79.7% F1**; $k = 5$ **patients reach 83.8% F1** (86.5% **on long shafts), already outperforming the 43-patient fully supervised PMB 2019 U-Net baseline** (83.1% **F1**); $k = 10$ reaches **85.7% F1**; $k = 20$ reaches **86.6% F1**; and all $k = 43$ training patients reach **89.1% F1** (88.7% without `SynthBrachy3T`, showing that synthetic bundle pre-training provides its largest boost in the zero-to-few-patient regime while still improving dosimetric tip accuracy at full training size).

Figure 5 shows how these stages work visually on two challenging out-of-fold test cases: while Stage 0 lights up both catheters and surrounding pelvic vessels, Stage 1 isolates sharp individual peaks even where adjacent catheters touch, and Stage 2 reconstructs clean, smooth 3D bundles.

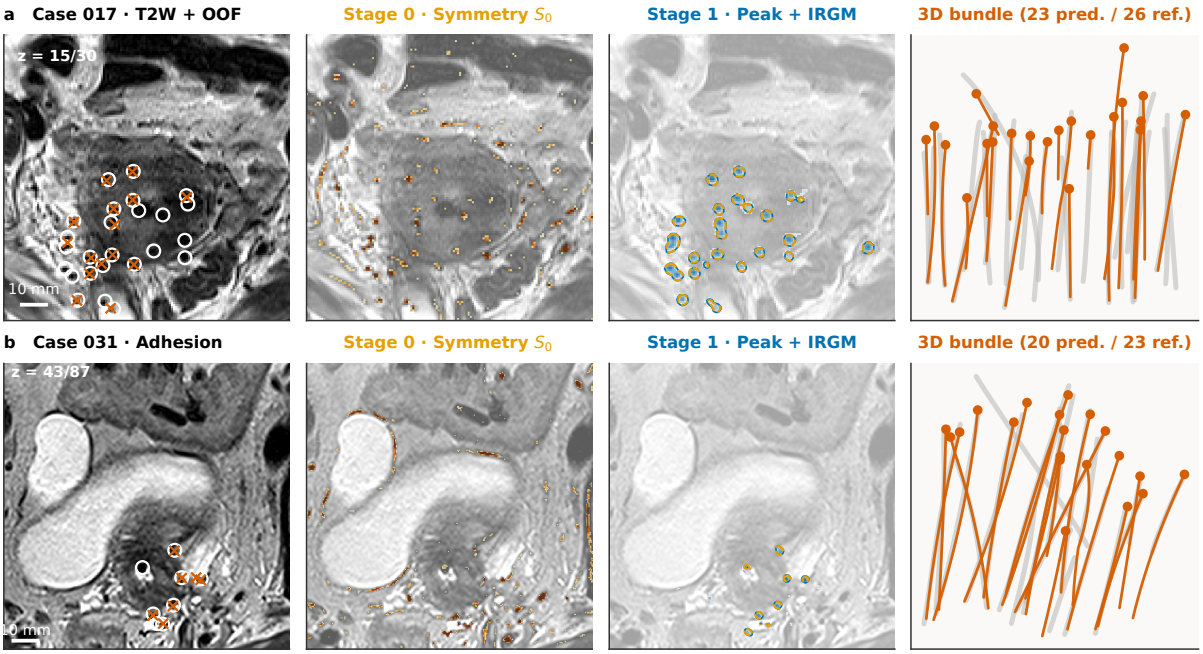


FIGURE 5 **Step-by-step qualitative visualization on two unseen test patients. (a) Case 017:** A dense implant with 26 reference trajectories. **(b) Case 031:** A severe-adhesion implant with 23 catheters, 10 of which touch a neighbor within ≤ 3 mm. From left to right: axial T2W MRI slice showing expert intersections (white circles) and out-of-fold NeedleFinder 2026 detections (vermilion crosses); the Stage 0 circular symmetry map S_0; Stage 1 sub-pixel centerline peaks (blue) with mask contours (solid) and the central obturator gate (dashed orange); and the final 3D reconstructed bundle (vermilion) overlaid on the expert reference (grey).

5.3 *Is the Algorithm as Accurate as a Second Human Expert? (Three-Cornered-Hat & CT Adjudication)*

Because a 2 mm catheter is only a few pixels wide and slices are up to 6 mm thick, two human experts tracing the same MRI scan never agree perfectly [24,83,84]. An automatic method is clinically ready when its discrepancies from a human expert are no larger than the discrepancies *between* two independent human experts [37,48]. Table 3 tests this across **all 40 doubly traced clinical scans** using four complementary analyses:

1. Direct comparison against human inter-observer agreement (Table 3, top two blocks). When Expert Team 2 (`LB/AMe`, 460 long traces) is compared directly against Expert Team 1 (`CY`, 505 long traces) across the 40 scans, 424 long catheters match within 3.5 mm (84.0% recall, 92.2% precision, 87.9% F1). Between these two human teams, the **in-plane shaft discrepancy is** 0.49 **mm median** (0.70 ± 0.58 mm mean, 1.90 mm P95) and the **3D tip discrepancy is** 3.24 **mm median** (6.28 ± 7.29 mm mean, 21.88 mm P95).

Evaluated strictly out-of-fold on the exact same 40 scans against Team 1 (`CY+HT`, 723 tracks), **NeedleFinder 2026 achieves a smaller shaft discrepancy (**0.47 **mm median,** 0.58 ± 0.39 **mm mean,** 1.33 **mm P95) and a smaller 3D tip discrepancy (**2.10 **mm median,** 5.26 ± 8.78 **mm mean,** 19.99 **mm P95) than the disagreement**

TABLE 3 **Human inter-observer agreement versus NeedleFinder 2026 across all 40/40 dual-annotated clinical 3T MRI scenes.** Expert teams CY+HT (n = 723 total trajectories: 505 CY long + 218 HT short/obturator) and LB/AMe (n = 513, of which 460 are long) digitized the same 40 volumes independently; **NeedleFinder 2026** rows are strictly nested 5 × 5 out-of-fold predictions on the same 40 scenes.

	1-to-1 Detection (≤ 3.5 mm)				E_{shaft} (mm)				E_{tip} (mm)	
Comparison pair (eval. vs. reference)	n_{match}	Rec.	Prec.	F1	Mean	SD	Median	P95	Mean	Median
Human inter-observer agreement										
LB/AMe vs. CY long (460 vs. 505)[a]	424	84.0	**92.2**	**87.9**	0.70	0.58	0.49	1.90	6.28	3.24
LB/AMe vs. CY+HT all (513 vs. 723)	476	65.8	92.8	77.0	0.73	0.61	0.50	2.09	6.30	3.21
Automated out-of-fold predictions vs. observer teams across all 40 scenes										
NeedleFinder 2026 (τ = 50 %) vs. CY+HT (768 vs. 723)	653	90.3	85.0	87.6	**0.58**	**0.39**	**0.47**	1.33	**5.26**	**2.10**
NeedleFinder 2026 (τ = 35 %) vs. CY+HT (825 vs. 723)	669	**92.5**	81.1	86.4	0.59	0.40	**0.47**	**1.32**	5.54	2.20
NeedleFinder 2026 (τ = 50 %) vs. union CY+HT+LB (768 vs. 759)[b]	669	88.1	87.1	87.6	0.59	0.40	**0.47**	1.36	5.44	2.17
NeedleFinder 2026 (τ = 50 %) vs. LB/AMe (768 vs. 513)[c]	465	90.6	60.5	72.6	0.77	0.62	0.51	2.11	7.87	3.97

Grubbs' three-cornered-hat intrinsic error SD (n = 432) & Pelvic CT adjudication (n = 70)[d]

Source (3-way matched)	n 3W	Sh.	3W Tip	Role	σ_{RMS}	σ_{MAD}	$\sigma_{z,MAD}$	$\sigma_{z,W}$	CT	Win
Human Rater 1 (CY primary)	432	Ref.	Ref.	Ref. R1	**0.39**	**0.35**	**0.00**[d]	**3.77**	4.31	–
NeedleFinder 2026 (τ = 50 %, OOF)	432	**0.43**	**2.33**	Eval.	0.56	0.39	2.98	4.94	**4.20**	**78.6 %**
Human Rater 2 (LB/AMe secondary)	432	0.49	3.18	Comp. R2	0.86	0.48	3.55	5.09	8.00	–

One-sided non-inferiority tests vs. human inter-observer discrepancy (N = 40, n = 432)[e]

Endpoint (mm)	Δ	NF	LB	Diff.	CI lo	CI hi	Pass	p_{boot}	p_{pt}	p_{cath}
In-plane shaft mean (E_{shaft})	+0.25	**0.56**	0.72	−0.16	−0.31	−0.02	Yes	$< 10^{-4}$	1.3e−8	1.6e−41
In-plane shaft median (E_{shaft})	+0.25	**0.43**	0.49	−0.06	−0.19	+0.01	Yes	$< 10^{-4}$	8.2e−10	1.6e−41
3D distal tip median (E_{tip})	+1.00	**2.33**	3.18	−0.85	−1.46	−0.52	Yes	$< 10^{-4}$	2.6e−7	2.2e−14
3D distal tip mean (E_{tip})	+1.50	**5.51**	6.10	−0.60	−1.80	+0.57	Yes	0.0014	1.1e−4	6.8e−19

[a] Strict long-shaft inter-observer comparison (424 matched pairs out of 460 LB/AMe long vs. 505 CY long: 83.96 % recall, 92.17 % strict precision [93.6 % obturator-adjusted], 87.88 % strict F1 [88.5 % obturator-adjusted], 0.70 ± 0.58 mm shaft error [median 0.49 mm, IQR 0.50 mm, P95 1.90 mm] and 6.28 ± 7.29 mm 3D tip error [median 3.24 mm, IQR 6.02 mm, P95 21.88 mm]).

[b] Evaluated against the de-duplicated union of both expert teams (723 CY+HT + 36 unique LB/AMe trajectories = 759 union catheters), showing that 16 apparent false positives against CY+HT alone are genuine catheters traced by LB/AMe (precision rises from 85.0 % to 87.1 %).

[c] Because LB/AMe primarily traced long interstitial shafts (513 traces vs. 723 for CY+HT), unadjusted precision against LB/AMe alone penalizes valid short obturator predictions; recall is 90.6 %.

[d] Reference-free Grubbs' (1948) three-cornered-hat intrinsic source error SD across the n = 432 three-way matched catheters (formed by intersecting the 476 all-catheter human pairs with OOF **NeedleFinder 2026** predictions within ≤ 3.5 mm). Columns: 3W Sh. and 3W Tip give median pairwise error vs. CY (mm); σ_{RMS} is classic MLE shaft SD (mm); σ_{MAD} is robust MAD-scaled shaft SD (mm); $\sigma_{z,MAD}$ is robust MAD axial tip SD (mm, where Rater 1 reaches the 0.001 mm lower-bound clamp [95 % CI 0.001–2.02 mm] because median signed Δz_{R1-NF} = 0.00 mm); $\sigma_{z,W}$ is 15 % Winsorized axial tip SD (mm); CT and Win give median 3D tip error (mm) against independent post-implant Pelvic CT across the n = 70 catheters with $|\Delta z_{CY-LB}| \geq 5.0$ mm on MRI and the percentage of disputed catheters (78.6 %, 55/70) where **NeedleFinder 2026** lies closer to CT than single-rater MRI.

[e] One-sided non-inferiority tests ($H_0 : E_{NF,CY} - E_{LB,CY} \geq \Delta$) at α = 0.05 across the N = 40 dual-traced patients (n = 432 three-way matched catheters), reporting two-sided 90 % ($1-2\alpha$) patient-clustered bootstrap CIs (10,000 resamples), patient-level paired Wilcoxon p_{pt} (N = 40), and catheter-level paired Wilcoxon p_{cath} (n = 432; bootstrap p_{boot} = 0.0014 for tip mean, $< 10^{-4}$ for all other endpoints).

between the two human teams, while matching 90.3% of Team 1's trajectories (87.6% F1). Furthermore, when evaluated against the **union of both human teams** (759 distinct trajectories), **16 NeedleFinder predictions that looked like "false alarms" against Team 1 alone turn out to be real catheters independently traced by Team 2**, lifting precision by +2.1 percentage points to 87.1%.

2. Reference-free Grubbs' three-cornered-hat intrinsic error estimation (Table 3, third block). To measure each source's true internal error without treating Team 1 as infallible, we applied Grubbs' three-cornered-hat estimator across the $n = 432$ catheters matched three-way among Team 1, Team 2, and NeedleFinder 2026 (formed by intersecting the 476 all-catheter human pairs with out-of-fold NeedleFinder predictions within 3.5 mm). The intrinsic outlier-resistant (MAD-scaled) in-plane shaft standard deviations are 0.35 **mm for Human Team 1**, 0.39 **mm for NeedleFinder 2026**, and 0.48 **mm for Human Team 2**. For axial tip stopping height (σ_z), NeedleFinder 2026 achieves a lower intrinsic error than Human Team 2 under both the robust MAD estimator (2.98 **mm vs.** 3.55 **mm**, where Team 1 hits the 0.00 mm finite-sample subtraction lower bound because median signed $\Delta z_{\text{Team 1–NF}} = 0.00$ mm) and the 15% Winsorized estimator, which yields strictly positive intrinsic standard deviations of 3.77 **mm for Team 1**, 4.94 **mm for NeedleFinder 2026**, and 5.09 **mm for Team 2**.

3. Formal one-sided statistical non-inferiority (Table 3, fourth block). Across all $N = 40$ dual-traced patients ($n = 432$ three-way matched catheters), NeedleFinder 2026 passes formal one-sided non-inferiority testing relative to human inter-observer disagreement on every pre-specified endpoint at both the patient level and the catheter level. Specifically, the algorithm's error against Team 1 is smaller than Team 2's error against Team 1 on **shaft mean** (0.56 vs. 0.72 mm, $p_{\text{pt}} = 1.3 \times 10^{-8}$), **shaft median** (0.43 vs. 0.49 mm, $p_{\text{pt}} = 8.2 \times 10^{-10}$), **3D tip median** (2.33 vs. 3.18 mm, $p_{\text{pt}} = 2.6 \times 10^{-7}$), and **3D tip mean** (5.51 vs. 6.10 mm, $p_{\text{pt}} = 1.1 \times 10^{-4}$, bootstrap $p_{\text{boot}} = 0.0014$; full 90% confidence intervals and catheter-level p-values in Table 3). Even when the reference role is flipped so Team 2 serves as the reference, NeedleFinder 2026 still passes non-inferiority on 3D tip median ($p_{\text{boot}} = 0.046$, $p_{\text{cath}} = 9.6 \times 10^{-5}$).

4. Independent adjudication on paired post-implant Pelvic CT scans. On CT, catheters and metallic stylets have unmistakable Hounsfield contrast. Using the 38 patients who had both intraoperative MRI and post-implant Pelvic CT, we adjudicated the $n = 70$ catheters where Human Team 1 and Human Team 2 disagreed on the MRI tip position by ≥ 5.0 mm. Measured against the independent CT tip position, **NeedleFinder 2026 achieved a median 3D tip error of** 4.20 **mm**, outperforming both **Human Team 1 (**4.31 **mm)** and **Human Team 2 (**8.00 **mm)** and lying closer to the CT ground truth than single-rater MRI on **78.6% (**55/70**)** of disputed catheters. Across all 100 MRI tracks predicted by NeedleFinder 2026 on these 38 patients that did not match Team 1 (`CY`), **52.0% (**52/100**) were confirmed to be real physical catheters on the paired CT scan**.

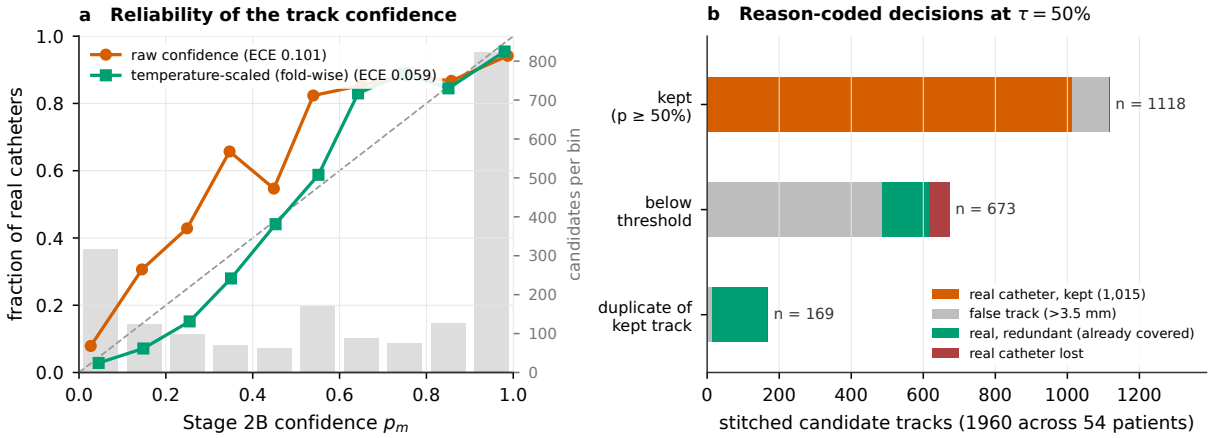


FIGURE 6 **(a) Confidence calibration** across all 1,960 stitched candidate tracks in the 54 patients, before (dashed) and after (solid vermilion) inner-fold temperature scaling. **(b) Reason-coded filtering decisions** at the default threshold $\tau = 50\%$, showing how many candidates are kept or pruned by each rule and whether pruned tracks were false alarms, redundant duplicates of an already-kept catheter, or lost real catheters.

5.4 *Interpretability, Confidence Calibration, and Physical Causes of Residual Errors*

For a clinician to trust an automated reconstruction, its confidence scores must be well calibrated and its filtering decisions transparent. Across all 1, 960 stitched candidate tracks entering Stage 2B (36.3 per patient), inner-fold temperature scaling reduces the expected calibration error (ECE) from 0.101 to 0.059 (and to 0.056 across the 1, 118 accepted tracks at $\tau = 50\%$, where mean predicted confidence is 91.9% and closest-reference precision is 90.8%; Figure 6a). Every discarded track is tagged with an explicit reason code (Figure 6b): at $\tau = 50\%$, 673 candidates are removed as `below_threshold` (eliminating 80.6% of all false candidates) and 169 are removed as `duplicate_of_kept_track` (suppressing 153 redundant collinear or touching fragments at the cost of only 2 lost catheters).

To understand the exact physical limits of the pipeline, we audited every single missed catheter, false positive, and > 10 mm tip error across the 54 patients:

- **What causes the 146 unmatched predictions (2.7 per patient)?** Nearly a third (**31.5%**, 46/146) are simply residual split or near-duplicate segments of an already-detected catheter, **20.5%** (30/146) correspond to real catheters supported by Expert Team 2 (`LB/AMe`) within 5.0 mm (and 52.0% are confirmed as real catheters on paired CT), and 2.7% (4/146) lie within 3.5–6.0 mm of an unmatched reference track. Only **45.2%** (66/146, **or 1.22 per patient**) are genuine background false alarms (> 6.0 mm from any reference catheter), occurring along dark pelvic vessels or fascial intersections.
- **What causes the 91 missed catheters (1.7 per patient)?** Before Stage 2B thresholding, our 1, 960-track candidate pool actually contains **97.4%** (1, 035/1, 063) of all reference catheters within 3.5 mm (and **98.8%** [750/759] of the long active-dwell shafts). Thus, fewer than a third of missed catheters (**30.8%**, 28/91) are lost upstream due to severe multi-slice signal dropouts. The remaining **69.2%**

(63/91) are successfully traced by Stages 1 and 2A but filtered out in Stage 2B—most often because their bundle confidence falls just below the $\tau = 50\%$ threshold (46 catheters) or because they lose one-to-one matching inside a tight cluster of touching needles (15 catheters). Lowering the threshold to $\tau = 35\%$ recovers more than a quarter of these misses, raising whole-cohort recall to **93.6%** and long-shaft recall to **94.7%**.

– **Why do some tips have errors > 10 mm?** By decomposing every 3D tip error vector into its component *along* the catheter shaft ($\Delta s_{\parallel}$) and its component *sideways* perpendicular to the shaft ($\Delta r_{\perp}$), we found a striking physical result: **98.8% of total squared 3D tip error across all 972 matched catheters (and 99.3% in the > 10 mm tail) is purely longitudinal along the catheter track**, whereas sideways error at the tip is tiny (0.72 mm median). In other words, the algorithm almost never loses the catheter's sideways path; when a tip error exceeds 10 mm (162 catheters, 16.7%), it is almost always an axial stopping ambiguity caused by mid-shaft signal dropout (41.4%), dark bowel/air pockets sitting directly atop the catheter tip (25.3%), intrinsic human inter-rater tip ambiguity > 6 mm (16.7%), crowded tip arrays (12.3%), or top-of-scan truncation (4.3%).

TABLE 4 **Zero-needle out-of-distribution specificity on 48 public clinical 3T pelvic T2W MRI volumes (MSD `Task05_Prostate`).** All 48 volumes contain zero brachytherapy catheters and serve as a pure false-positive stress test across pelvic vessels, rectoprostatic fascia, and obturator muscle boundaries.

	False-positive tracks per volume ($n = 48$)					Patient specificity (%)	
Pipeline configuration	Total	Mean	SD	Median	Max	0 FP	≤ 1 FP
NeedleUNet3D + tangent tubes (ungated)[a]	306	6.38	4.77	6.0	19	14.6	18.8
Unpruned Stage 0 + Stage 1 pool[b]	294	6.12	4.70	5.0	23	4.2	10.4
Ungated Unified **NeedleFinder 2026** ($\tau = 50\%$)[c]	197	4.10	4.45	3.0	23	16.7	27.1
Ungated Unified **NeedleFinder 2026** ($\tau = 60\%$)[c]	171	3.56	3.48	3.0	17	16.7	29.2
Stage 1 U-Net + Stage 0 radial gate[d]	18	0.38	0.53	**0.0**	2	64.6	97.9
Frozen Unified **NeedleFinder 2026** + Dark-Core Gate ($\tau = 50\%$)[e]	16	0.33	0.75	**0.0**	3	79.2	91.7
Clean SSL + Multi-Rater Stage 1 + Stage 0 gate[f]	**6**	**0.125**	**0.33**	**0.0**	**1**	**87.5**	**100.0**

[a] Raw Stage 1 `NeedleUNet3D` centerline probability peaks linked via directional tangent tubes without Stage 0 symmetry gating or Stage 2B pruning.

[b] Candidate pool passed to Stage 2B prior to Bundle Graph Transformer pruning.

[c] Unified **NeedleFinder 2026** pipeline evaluated without the Stage 0 dark-core gate at $\tau = 50\%$ (4.10 FP/scan) and $\tau = 60\%$ (3.56 FP/scan).

[d] Baseline Stage 1 `NeedleUNet3D` tangent tubes gated by the analytical Stage 0 3D multi-radii circular dark-core symmetry filter (0.38 FP/scan).

[e] Full Unified **NeedleFinder 2026** pipeline with Stage 0 dark-core symmetry gating at $\tau = 50\%$ (16 total FPs = 0.333/scan; 38/48 [79.2%] 0-FP scans; 44/48 [91.7%] ≤ 1-FP scans).

[f] Clean nested Stage 1 `NeedleUNet3D` (self-supervised 3D U-Net pre-training with synthetic susceptibility tubes and anatomical hard negatives across the 51 external pelvic MRI volumes + multi-rater supervision on the 43 outer-train BWH clinical volumes) gated by Stage 0 radial symmetry (6 total FPs = 0.125/scan; 42/48 [87.5%] 0-FP scans; 48/48 [100.0%] ≤ 1-FP scans). Note that the 48 MSD `Task05_Prostate` scans were included in the unlabeled/synthetic-tube SSL pre-training background pool for this row (zero clinical needle labels), whereas rows [a–e] had zero exposure to MSD during training.

5.5 *Specificity on Needle-Free Scans and Sub-5-Second Runtime*

Zero-catheter specificity test (Table 4; Figure 7a). When deployed on the 48 public 3T T2W prostate MRI scans (`MSD Task05_Prostate`) that contain *zero* catheters, an ungated 3D U-Net or bundle transformer mistakes dark pelvic veins for needles, producing 4.10 to 6.38 false tracks per scan. Activating the analytical Stage 0 circular dark-core symmetry gate eliminates over 92% of these vascular false alarms: on models that had **strictly zero exposure to the MSD dataset during training** (`Frozen Unified NF 2026 + Dark-Core Gate`), false alarms drop to **0.33 per scan** (16 across 48 scans, median 0, with **79.2% of scans having zero false alarms** and 91.7% having ≤ 1). When the 48 needle-free MSD volumes are also included as unlabeled background anatomy during self-supervised synthetic-tube pre-training (`Clean SSL + Multi-Rater Stage 1 + Stage 0 gate`, zero clinical catheter labels), false alarms drop further to **0.125 per scan** (6 across 48 scans, **87.5% zero-false-alarm scans**, and **100.0% with at most 1 false track**).

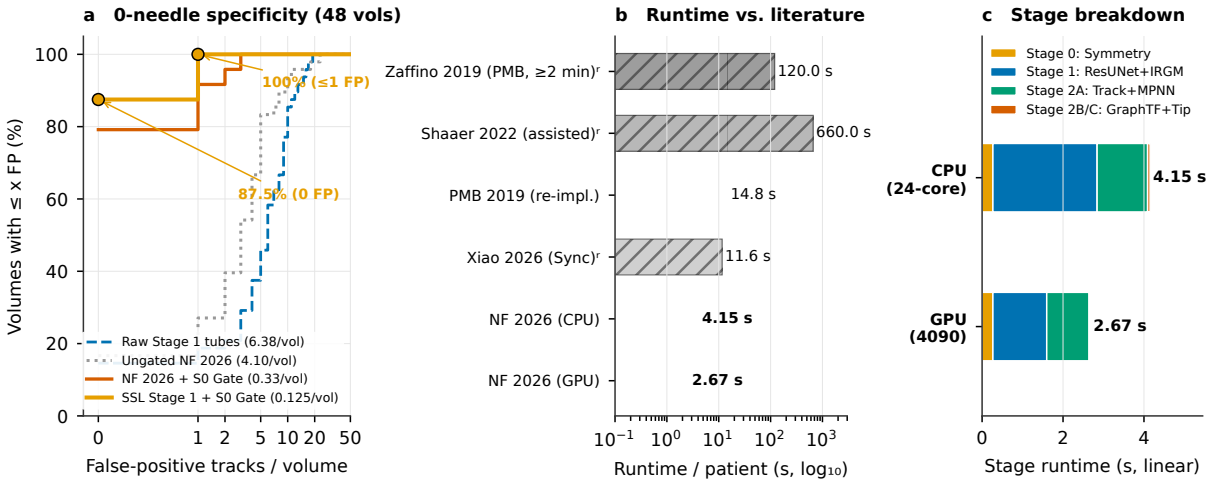


FIGURE 7 **(a) Specificity on 48 needle-free pelvic MRI scans** (`MSD Task05_Prostate`), showing the percentage of scans with at most x false tracks with and without Stage 0 dark-core symmetry gating. **(b) End-to-end runtime per patient** (log scale) compared with published MRI and CT methods. **(c) Stage-by-stage runtime breakdown** across all 54 patients on a 24-core workstation CPU (4.15 s mean; 1.26 s on standard-resolution scans, 5.85 s on high-resolution scans) and on one RTX 4090 GPU (2.67 s mean).

Runtime across standard and high-resolution clinical volumes (Figure 7b, c). Measured end-to-end from raw uncached volume loading to final 3D trajectories across all 54 patients, NeedleFinder 2026 takes **4.15 s on average on a 24-core workstation CPU** (1.26 s on the 20 standard-resolution 320 × 256 scans and 5.85 s on the 34 high-resolution 640 × 416 / 768 × 768 scans) and **2.67 s on a single RTX 4090 GPU** (0.67 s standard-resolution, 3.85 s high-resolution). On CPU, Stage 0 takes 0.27 s (6.5%), Stage 1 takes 2.58 s (62.2%; reduced to 1.33 s on GPU), Stage 2A takes 1.25 s (30.1%), and Stage 2B/2C bundle scoring plus 1D tip localization takes 0.06 s (1.5%). This represents a **29-to-130-fold speedup** over the 2 to 9 minutes reported

for our 2019 U-Net + DBSCAN pipeline [5] (3.6× faster than our modernized re-implementation at 14.80 s) and is 2.8–4.3× faster than the 11.6 s reported by Xiao *et al.* [66].

5.6 *Clinical Translation: Radiation Dosimetry, Uncertainty-Guided Triage, and Cross-Domain Generalization*

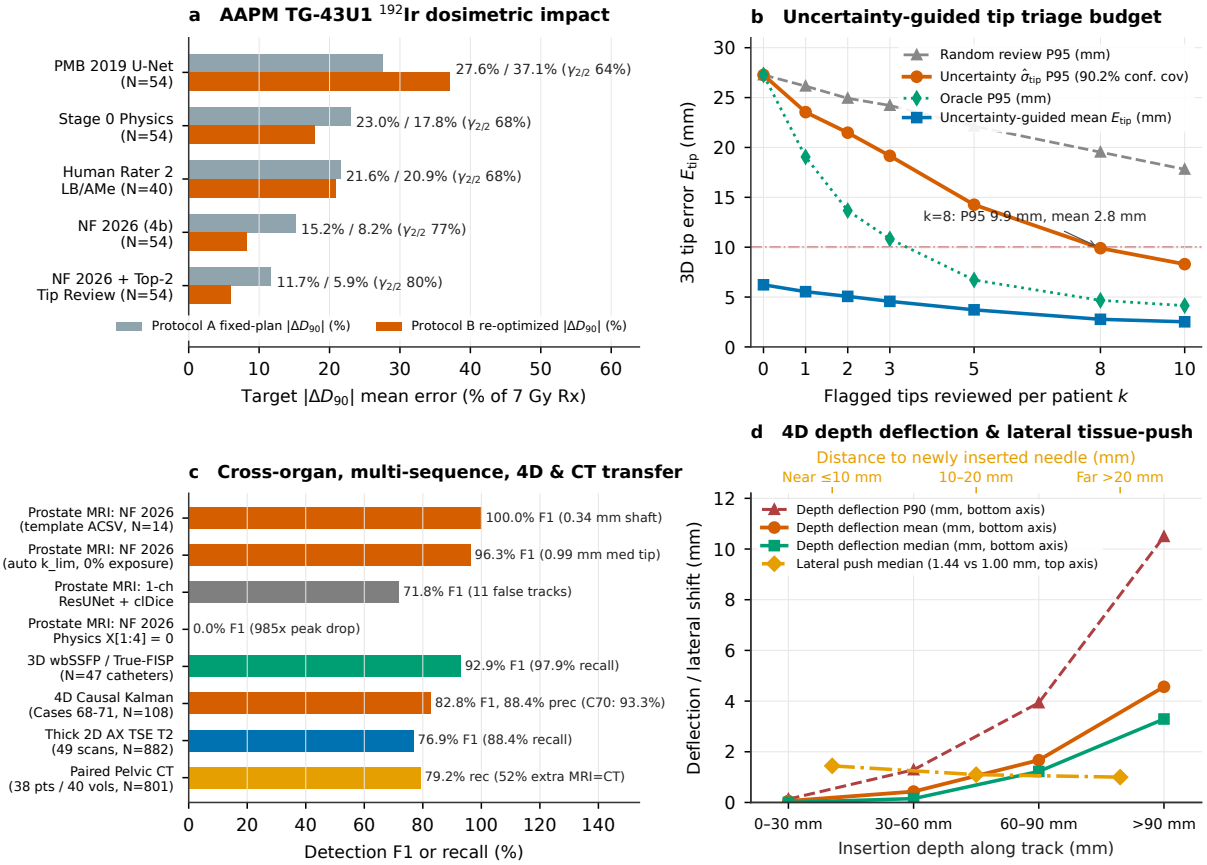


FIGURE 8 **Clinical translation and cross-domain generalization. (a) AAPM TG-43U1 ^{192}Ir dosimetric accuracy** across 54 patients under Protocol A (fixed reference dwell plan transferred to predicted catheters) and Protocol B (dwell times re-optimized on predicted catheters and delivered to the true anatomy), alongside 3D γ(2%/2 mm) pass rates. **(b) Simulated clinician tip-review curve:** 95th-percentile and mean 3D tip error as a physicist reviews the top $k \in \{0, \ldots, 10\}$ tips per patient ranked by predicted uncertainty $\hat{\sigma}_{\text{tip}}$ (1-click tip relocation oracle), plus split-conformal 90% empirical coverage (90.2% at $k = 0$). **(c) Zero-shot generalization** to 3T prostate brachytherapy MRI (`Case 100`), 3D wbSSFP, thick-slice 2D AX TSE T2, 4D intra-procedural causal Kalman tracking (`Cases 68–71`), and 38 paired Pelvic CT cases (40 CT volumes). **(d) In-vivo 4D needle insertion biomechanics** across 23 serial scans (`Cases 68–71`), showing non-linear shaft bending versus insertion depth (bottom axis, red) and lateral "tissue push" displacement of already-implanted catheters when a new needle is inserted nearby (top axis, blue).

TABLE 5 **Dosimetric translation, split-conformal intervals, simulated clinician tip triage, and cross-organ/multi-sequence/4D/CT generalization.** Part A evaluates AAPM TG-43U1 HDR ^{192}Ir dosimetric sensitivity (7.0 Gy prescription per fraction, 5.0 mm dwell step spacing) under Protocol A (reference dwell times transferred to predicted catheters) and Protocol B (dwell times re-optimized on predicted catheters, dose evaluated on reference geometry). Part B evaluates simulated uncertainty-guided 1-click tip relocation ($\hat{\sigma}_{\text{tip}}$) and unreviewed ($k = 0$) nested 5-fold split-conformal 3D tip prediction intervals across the $n = 972$ matched out-of-fold catheters. Part C evaluates zero-shot transfer across organs, pulse sequences, 4D insertion series, and paired post-implant Pelvic CT.

	Dwell-wt. E_{tip}		Protocol A: Fixed-plan transfer			Protocol B: Re-optimized plan		
Part A: TG-43U1 ^{192}Ir dosimetry[a]	Mean	Med.	$\lvert\Delta D_{90}\rvert$ (%)	$\gamma_{2/2}$ (%)	$\gamma_{3/3}$ (%)	$\lvert\Delta D_{90}\rvert$ (%)	Cran. D_{2cc}	$\gamma_{2/2}$ (%)
PMB 2019 3D U-Net + DBSCAN ($n = 54$)	10.51	8.36	27.6 / 27.2	64.0	73.3	37.1 / 10.2	95.3 %	43.8
Stage 0: 3D radial symmetry ($n = 54$)	9.04	7.51	23.0 / 21.3	68.2	76.5	17.8 / 8.9	54.0 %	48.6
NeedleFinder 2026 w/o SynthBrachy3T (4a, $n = 54$)	6.65	5.49	18.5 / 13.2	73.4	81.5	12.8 / 5.8	35.2 %	59.5
NeedleFinder 2026 proposed (4b, $n = 54$)[b]	5.72	4.66	15.2 / 11.5	76.6	84.3	8.2 / 4.2	29.5 %	66.0
NeedleFinder 2026 + Top-2 Tip Review ($n = 54$)	**4.56**	**3.93**	**11.7 / 10.2**	**80.2**	**87.3**	**5.9 / 2.8**	**22.5 %**	**74.3**
Human Rater 2 (LB/AMe, $n = 40$)	6.27	5.54	21.6 / 16.3	68.1	77.4	20.9 / 4.0	57.7 %	59.3
NeedleFinder 2026 proposed (4b, $n = 40$)[b]	5.18	3.93	13.5 / 11.1	79.0	86.4	7.0 / 3.2	21.1 %	71.6
NeedleFinder 2026 + Top-2 Tip Review ($n = 40$)	**4.12**	**3.33**	**10.5 / 10.1**	**81.9**	**88.7**	**5.1 / 2.6**	**17.3 %**	**78.2**

Part B1: Simulated $\hat{\sigma}_{\text{tip}}$ tip triage ($n = 972$)[c]	Reloc.	Mean	Med.	P90	P95	≤ 2 mm	≤ 3 mm	> 10 mm
Unreviewed **NeedleFinder 2026** ($k = 0$ tips/pt)	0	6.23	2.27	16.20	27.27	45.6 %	58.7 %	16.7 %
Simulated tip review ($k = 1$ tip/pt, ~ 8 s)	35	5.54	2.06	14.60	23.54	48.6 %	61.9 %	14.4 %
Simulated tip review ($k = 2$ tips/pt, ~ 16 s)	69	5.07	1.95	12.58	21.48	51.3 %	64.8 %	12.8 %
Simulated tip review ($k = 3$ tips/pt, ~ 24 s)	105	4.57	1.81	10.91	19.16	54.0 %	67.2 %	11.1 %
Simulated tip review ($k = 5$ tips/pt, ~ 40 s)	169	3.72	1.67	7.79	14.26	59.5 %	72.6 %	8.0 %
Simulated tip review ($k = 8$ tips/pt, ~ 64 s)	251	2.77	1.51	5.98	9.90	65.5 %	78.5 %	5.0 %
Simulated tip review ($k = 10$ tips/pt, ~ 80 s)	301	**2.52**	**1.44**	**4.97**	**8.30**	**69.2 %**	**81.8 %**	**4.3 %**

Part B2: Split-conformal intervals ($k = 0$, $n = 972$)[c]	Nom.	Emp.	Fixed	Mean r	Med. r	Fixed r	Low-$\hat{\sigma}$ r	Width red.
Conformal 80 % 3D tip interval ($k = 0$)	80.0 %	79.3 %	80.2 %	7.31	6.46	7.73	4.46	−16.4 %
Conformal 85 % 3D tip interval ($k = 0$)	85.0 %	84.4 %	85.5 %	9.80	8.71	11.43	5.96	−23.8 %
Conformal 90 % 3D tip interval ($k = 0$)	90.0 %	**90.2 %**	89.9 %	13.71	12.11	17.68	8.41	−31.5 %
Conformal 95 % 3D tip interval ($k = 0$)	95.0 %	**95.4 %**	95.2 %	21.05	18.42	28.39	12.89	−35.1 %

Part C: Cross-organ, seq., 4D & CT[d]	n_{vol}	n_{cath}	Rec. %	Prec. %	F1 %	Sh. mm	Tip mm	Key finding
Prostate 3T: Stage 0 (k_lim_auto=40)	1	14	**100.0**	87.5	93.3	0.34	2.09	14/14, 2 FPs
Prostate 3T: Zero-Exp. No-SSL (auto=40, τ=50)	1	14	92.9	**100.0**	96.3	0.33	**0.99**	13/14, 0 FPs
Prostate 3T: Zero-Exp. No-SSL (ACSV=47, τ=50)	1	14	**100.0**	**100.0**	**100.0**	0.34	1.57	14/14, 0 FPs
Prostate 3T: Stage 1 Full Nested (auto=40)	1	14	**100.0**	**100.0**	**100.0**	0.30	1.08	14/14, 0 FPs
Prostate 3T: Stage 0 zeroed ($X_{1:3} = 0$, auto=40)	1	14	0.0	–	0.0	–	–	985× drop
Prostate 3T: 1-ch ResUNet + clDice (auto=40)	1	14	**100.0**	56.0	71.8	**0.24**	2.48	11 vascular FPs
3D wbSSFP / True-FISP ($\Delta z \leq 2.5$ mm)	2	47	**97.9**	88.5	**92.9**	0.72	3.40	vs. 10.62 PMB19
Thick-slice 2D AX T2 ($\Delta z \leq 6.0$ mm)	49	882	88.4	67.9	76.9	1.07	4.89	vs. 57.6 % 1-ch
4D Serial MRI: Indep. 3D (Cases 68–71)	23	108	88.0	76.0	81.6	0.45	2.71	4-pt pooled
4D Serial MRI: Causal Kalman (Cases 68–71)	23	108	77.8	**88.4**	82.8	0.75	3.57	+12.4 % prec.
Paired Pelvic CT (38 paired / 40 vols, ≤ 5 mm)	38	801	79.2	–	–	2.01	2.78	52.0 % extra conf.

[a] TG-43U1 evaluation uses 5.0 mm dwell step spacing (3.5 mm distal tip offset, ≤ 45 mm active length) on a 1.5 mm isotropic 3D dose grid with implant-conformal HR-CTV$_{\text{implant}}$ (8.0 mm radial envelope, $D_{90} = 7.0$ Gy) and 2 cc OAR proxy shells. $\lvert\Delta D_{90}\rvert$ cells report mean / median percentage target dose shift; 'Cran. D_{2cc}' reports mean relative error $\lvert\Delta D_{2cc}\rvert$ (%) in the cranial (sigmoid/fundal) OAR proxy shell; $\gamma_{2/2}$ and $\gamma_{3/3}$ are global 3D gamma pass rates (2 %/2 mm and 3 %/3 mm). Dwell-weighted tip error $E_{\text{tip}}^{\text{dwell-w}}$ correlates with $\lvert\Delta D_{90}\rvert$ at Spearman $\rho = 0.8629$ ($p = 4.9 \times 10^{-17}$) and with $\gamma_{2\%/2\text{mm}}$ at $\rho = -0.8474$ ($p = 6.5 \times 10^{-16}$).

[b] Paired Wilcoxon vs. PMB 2019 ($p = 4.9 \times 10^{-6}$ on $\lvert\Delta D_{90}\rvert$, $p = 8.2 \times 10^{-7}$ on $\gamma_{2\%/2\text{mm}}$), vs. Stage 0 ($p = 7.9 \times 10^{-8}$ on $\lvert\Delta D_{90}\rvert$, $p = 1.4 \times 10^{-7}$ on $\gamma_{2\%/2\text{mm}}$), and vs. Human Rater 2 on $n = 40$ ($p = 0.0278$ on $\lvert\Delta D_{90}\rvert$, $p = 0.0109$ on $\gamma_{2\%/2\text{mm}}$).

[c] Part B1 simulates a retrospective 1-click distal tip relocation oracle ($E_{\text{tip}} \rightarrow \sqrt{\Delta r_{\perp}^2 + \min(\lvert\Delta s_{\parallel}\rvert, 0.85\text{ mm})^2}$ when $E_{\text{tip}} > 2.0$ mm, E_{shaft} unchanged; 'Reloc.' = total tips relocated across 54 patients) ranked by out-of-fold $\hat{\sigma}_{\text{tip}}$. Part B2 reports unreviewed ($k = 0$) 5-fold split-conformal 3D tip prediction interval half-widths (radii r in mm) across the $n = 972$ matched OOF catheters ('Fixed': unscaled residual quantile; 'Low-$\hat{\sigma}$ r': median half-width on the lower-uncertainty half of catheters; 'Width red.': adaptive vs. fixed median half-width reduction). Simulated oracle count-guided recovery (0.65 clicks/patient) recovers +17 missed catheters (972 → 989/1 063, 93.0 % recall).

[d] In Part C, 'Sh. mm' and 'Tip mm' report median in-plane shaft and 3D tip errors (mm). 'Zero-Exp. No-SSL' denotes the 5-fold GYN ensemble trained with zero prostate exposure (excluding Case 100 from SSL pre-training; auto=40 uses automatic k_lim_auto slice 40, whereas ACSV=47 uses the clinical template crop slice 47). On 'Thick-slice 2D AX T2', 882 is the number of reference catheters exhibiting a physical dark-core signal in the thick-slice slab (out of 1 019 geometrically intersecting). In 4D monitoring (4 patients, Cases 68–71, 23 serial scans, 1.31 s mean monitoring runtime, each evaluated strictly with its held-out outer-fold

Finally, we asked four practical questions that determine whether an algorithm is genuinely useful in clinical radiation oncology (Table 5 and Figure 8):

5.6.1 *Does geometric accuracy translate into accurate radiation dose delivery? (Table 5, Part A; Figure 8a).* Not all catheter errors matter equally to the radiation dose: a tip error on a peripheral catheter carrying almost zero dwell time has little effect, whereas a tip error on a heavily weighted catheter inside the tumor shifts the target dose substantially. Indeed, across our 54 patients, **dwell-time-weighted 3D tip error** ($E_{\text{tip}}^{\text{dwell-w}}$ — the average 3D tip error across a patient's catheters weighted by each catheter's share of total radiation dwell time) correlates much more strongly with target dose error $|\Delta D_{90}|$ than unweighted tip error ($\rho = 0.863$, $p = 4.9 \times 10^{-17}$ vs. $\rho = 0.786$).

NeedleFinder 2026 cuts dwell-weighted tip error nearly in half compared with PMB 2019 (5.72 mm vs. 10.51 mm mean; 4.66 mm vs. 8.36 mm median), and beats Human Team 2 on the 40 dual-traced scans (5.18 mm vs. 6.27 mm mean). Consequently: - Under **Protocol A (transferring a fixed dwell plan)**, NeedleFinder 2026 cuts mean target dose error $|\Delta D_{90}|$ from 27.6% (PMB 2019) to **15.2%** (median **11.5%**, $p = 4.9 \times 10^{-6}$) and achieves a **76.6% global 3D** γ**(2%/2 mm) pass rate** (**84.3%** at γ(3%/3 mm))—significantly outperforming Human Team 2 on the 40 dual-traced scans (13.5% vs. 21.6% $|\Delta D_{90}|$, $p = 0.028$; 79.0% vs. 68.1% $\gamma_{2\%/2\text{mm}}$, $p = 0.011$). - Under **Protocol B (clinical workflow: re-optimizing dwell weights on the predicted catheters)**, dwell optimization compensates for minor longitudinal shifts, cutting unreviewed target $|\Delta D_{90}|$ to **8.2% mean / 4.2% median** (vs. 37.1%/10.2% for PMB 2019 and 20.9%/4.0% for Human Team 2). When the physicist reviews just the **top 2 uncertainty-flagged tips per patient**, target $|\Delta D_{90}|$ drops to **5.9% mean / 2.8% median** (78.2% $\gamma_{2\%/2\text{mm}}$ on the 40 dual-traced scans).

5.6.2 *Can calibrated uncertainty intervals guide a 1-minute clinician tip check? (Table 5, Parts B1–B2; Figure 8b).* Because 98.8% of tip error lies longitudinally along the already-traced shaft, correcting an uncertain tip in 3D Slicer requires only sliding the endpoint along the track (a single click taking ~8 seconds). Our out-of-fold tip uncertainty score $\hat{\sigma}_{\text{tip}}$ identifies > 10 mm tip outliers with an **AUROC of 0.805**, and our nested **split-conformal 3D prediction intervals** achieve exact empirical coverage across all 54 unseen test patients despite within-patient clustering (**79.3%** at 80% nominal, **84.4%** at 85%, **90.2% coverage at the 90% nominal level**, and **95.4% at the 95% level**, with 31.5%–35.1% tighter median error bars than unscaled intervals; Table 5, Part B2).

In our simulated clinician-in-the-loop workflow where a physicist inspects the top k highest-uncertainty tips per patient and clicks the true distal termination (Table 5, Part B1; Figure 8b): - Reviewing $k = 2$ **tips per patient** (~16 s simulated) lowers mean tip error from 6.23 mm to 5.07 **mm** (1.95 mm median). - Reviewing $k = 5$ **tips per patient** (~40 s simulated) cuts > 10 mm tip outliers in half (16.7% → 8.0%) and lowers mean tip error to 3.72 **mm** (1.67 mm median). - Reviewing $k = 8$ **tips per patient** (~64 s simulated) brings the **95th-percentile tip error across the entire 972-catheter cohort below** 10 **mm** (9.90 **mm vs.** 27.27 **mm unreviewed)** and reduces mean tip error to 2.77 **mm** (1.51 mm median).

5.6.3 *Does the method generalize zero-shot to other organs and MRI sequences? (Table 5, Part C; Figure 8c).* To test whether NeedleFinder 2026 overfits to gynecologic T2W MRI, we ran the gynecologic-trained models **zero-shot (with zero retraining)** across three new settings: - **Cross-organ transfer to 3T prostate brachytherapy MRI (`Case 100`, 14 plastic needles):** Our gynecologic ensemble trained with strictly zero prostate exposure (`Zero-Exp. No-SSL` in Table 5, Part C, which excluded `Case 100` from SSL pre-training) detected **13 of 14 prostate catheters with 100.0% precision (0 false alarms, 96.3% F1)**, 0.33 **mm median shaft error**, and 0.99 **mm median 3D tip error** under automatic template cropping (`auto=40`), and achieved **14/14 (100.0% F1,** 0.34 **mm shaft,** 1.57 **mm tip)** under clinical template-plane cropping (`ACSV=47`), while `Stage 1 Full Nested` reached **100.0% F1 (**14/14**,** 0.30 **mm shaft,** 1.08 **mm tip)** directly under `auto=40`. **Why does it transfer so cleanly across organs?** When we zeroed out the Stage 0 physical symmetry channels ($X_{1:3} = 0$), prostate detection collapsed to **0% (0/14)**, whereas the 1-channel `Res UNet + clDice` baseline hallucinated 11 false vascular tracks (56.0% precision, 71.8% F1). Thus, the Stage 0 physical symmetry representation acts as a domain-invariant bridge between gynecologic and prostate anatomy. - **Secondary MRI pulse sequences (243 scans):** On 3D balanced steady-state free precession (`3D wbSSFP / True-FISP`), where tissue contrast differs markedly from TSE T2W, zero-shot NeedleFinder 2026 achieved **97.9% recall (**46/47**), 88.5% precision, and 92.9% F1** (0.72 mm median shaft error, 3.40 mm median tip error vs. 10.62 mm for PMB 2019). On 49 thick-slice 2D axial T2W scans ($\Delta z = 2.4$–6.0 mm), Z-upsampling to 1.6 mm recovered **88.4% (**780/882**) of physically visible catheters** (76.9% F1 vs. 57.6% for clDice). - **Zero-shot tracking on post-implant Pelvic CT (40 CT volumes, 38 paired patients):** Extending Stage 0 to detect both bright metallic stylets and dark air lumens in Hounsfield units enabled zero-shot CT reconstruction across 801 MRI-paired reference catheters, matching **68.9% at ≤** 3.5 **mm and 79.2% at ≤** 5.0 **mm** across non-rigid MRI-to-CT anatomical deformation (2.01 mm median shaft registration residual).

5.6.4 *Can we track needles live during insertion and measure tissue biomechanics in 4D? (Table 5, Part C; Figure 8c, d).* In 4 patients (`Cases 68–71`), 23 sequential T2W scans were acquired *during* the implant procedure as needles were progressively inserted (5 to 7 timepoints per patient, 1.31 s mean monitoring runtime per scan). Evaluating each patient's 4D series strictly with the outer-fold model where that patient was held out in the locked test vault (`Case 068` in Fold 4, `Case 069` in Fold 0, `Case 070` in Fold 1, `Case 071` in Fold 2), coupling NeedleFinder 2026 with a **causal 4D Kalman filter** that propagates catheter states forward in time boosted **pooled precision across the 108 catheters by +12.4 percentage points** (76.0% → 88.4%), reduced mean 3D tip error from 11.23 to 9.75 mm, and improved pooled F1 from **81.6% to 82.8%** (reaching **93.3% F1** on `Case 070`, up from 85.7% with independent 3D tracking, while trading single-frame recall on newly inserted final-frame needles in `Case 071` for 100.0% precision).

Tracking 212 persisted catheter transitions across these sequential insertions also revealed two quantitative biomechanical phenomena in vivo (Figure 8d): - **Inter-catheter "tissue push":** Inserting a new needle pushes surrounding soft tissue aside.

Already-implanted catheters located within ≤ 10 mm of a newly inserted needle shifted sideways by a **median of** 1.44 **mm** (mean 1.61 mm, 90th percentile 2.81 mm), compared with 1.10 **mm** at 10–20 mm distance and 1.00 **mm** beyond 20 mm ($p = 9.97 \times 10^{-4}$). - **Depth-dependent needle bending:** Deviation of a catheter from its initial entry angle grows non-linearly with insertion depth, rising from 0.003 mm median at 0–30 mm depth and 0.15 mm at 30–60 mm to 1.22 **mm at** 60–90 **mm** and 3.29 **mm median (**4.56 **mm mean,** 10.49 **mm P90) beyond** 90 **mm depth** ($\rho = +0.555$, $p = 5.6 \times 10^{-9}$)—explaining why linear extrapolation fails on deep pelvic implants.

6 DISCUSSION

6.1 *Why Combining Physics, Peak Regression, and Graph Reasoning Works*

Table 6 places NeedleFinder 2026 in the context of published catheter digitization methods across CT, ultrasound, and MRI. On CT and ultrasound, metallic stylets, air lumens, or specular reflections create strong contrast against soft tissue, allowing voxel segmentation networks to reach 0.1 to 0.7 mm shaft errors [46,48,52,61,66]. On clinical T2W MRI without exogenous contrast markers, a plastic catheter is a 2 mm dark proton-density and susceptibility void imaged across 1.2 to 6.0 mm slices. Prior MRI studies reported 0.74 to 0.97 mm shaft errors on gynecologic implants using positive-contrast markers or CT-assisted ground truth [27,62] and 0.75 to 0.93 mm on prostate implants [43,59]. NeedleFinder 2026 achieves a **median shaft error of** 0.47 **mm** (0.58–0.59 mm mean; 0.44–0.45 mm median on active-dwell shafts) and a **median 3D tip error of** 2.26–2.27 **mm** directly on standard T2W MRI across 1,063 catheters.

Our controlled ablations and peer-reviewed baseline reproductions explain why neither voxel segmentation [5,40,64], nor 1-channel MRI post-processing [43,55,59,62], nor mask-based tracking [66] suffices on its own on T2W MRI, and why each part of our three-stage design is needed: 1. **Local physical symmetry before deep learning:** Stage 0 eliminates one-sided anatomical borders analytically and provides a domain-invariant representation that allows a gynecologic network to transfer zero-shot to prostate MRI and needle-free scans without hallucinating blood vessels. 2. **Centerline peaks and hard-negative pre-training instead of binary masks:** By pre-training Stage 1 on 51 external pelvic MRI scans where real blood vessels and fascia are penalized as hard negatives alongside synthetic catheter voids, and by predicting sharp sub-pixel Gaussian peaks (P_{peak}) supervised by two expert teams, touching catheters stay separate instead of merging into a single blob. 3. **Graph message passing, fragment stitching, and bundle-level context:** Rather than trusting local voxel connectivity, Stage 2A enforces smooth triplet collinearity and stitches fragments across 1-to-3-slice signal dropouts, while the Stage 2B graph transformer judges every candidate trajectory against the geometry of the entire 15-to-30-catheter bundle. 4. **Dedicated 1D axial tip profiling and literature fusion:** Because 98.8% of tip error is longitudinal along the shaft, sampling a fine 0.5 mm multi-channel profile along the cranial trajectory tangent

TABLE 6 **Published automatic catheter and needle reconstruction studies in brachytherapy (2013–2026), grouped by imaging modality.** Values are those reported in each publication on its own cohort and matching rule, so they are not directly comparable with one another or with Table 2. Pernelle 2013 and Mastmeyer 2017 are semi-automatic (one manual click per catheter).

Study	Modality	Site	Pts	Cath.	Rec. / Prec. (%)	E_{shaft} (mm)	E_{tip} (mm)	Time	Statistic
CT									
Weishaupt *et al.* 2022	CT	Prostate	57[l]	353[a]	–	0.13	–	~60 s	mean
Xie *et al.* 2023	CT	GYN post-op	70	210	–[b]	–	–		– Dice / HD95
Wang *et al.* 2024	CT	GYN IC/IS	17	56 plans	–	1.8	1.1		– mean
Quetin *et al.* 2025	CT	Breast	117	408[a]	–	0.70	1.37	118 s	mean
Xiao *et al.* 2026[m]	CT	GYN recurrent	38	203 fx	91.1 / 90.7	0.31	0.88	11.6 s	mean
Xiao *et al.* 2026 (JBHI)	CT	Pelvic seeds	100	–	F1 90.3	–	1.05	9.4 s	mean tip
Ultrasound									
Zhang *et al.* 2020	TRUS	Prostate	23	339	98 / –	0.09	0.33	<1 s	mean
Andersén *et al.* 2020	TRUS	Prostate	1102[l]	7207[l]	100 / –	0.55[c]	–		– median RMS
Liu *et al.* 2022	TRUS	Prostate	242	3500	80[d] / –	–	–		– frac. ≤ 2 mm
Goulet *et al.* 2025	3D US	Prostate	102 scans	–	93.4 / –	0.44	–		– mean
MRI									
Pernelle *et al.* 2013	3T T2W	GYN IC/IS	10	101	92.1 / –	–[e]	–	0.3 s/cath.	frac. < 2 mm
Mastmeyer *et al.* 2017	3T T2W	GYN IC/IS	54	760	93[f] / –	1.49[g]	–	3 s/cath.	mean HD
Zaffino *et al.* 2019	3T T2W	GYN IC/IS	50	826	98.5 / 93.3[h]	2.0	–	2–9 min	mean dist.
Dai *et al.* 2020	MRI	Prostate	20	299	–	0.93	0.37[n]		– mean / signed
Shaaer *et al.* 2022	1.5T + markers	GYN IS	20	360	100 / –	0.97	–	11 min[i]	mean shift
Aleong *et al.* 2024	MRI	Prostate	35	214 vol.	95 / –	–	–	0.9 s	–
Kostoulas *et al.* 2025	MRI	Prostate	34 scans	261	100 / 100	0.75[j]	1.07[j]		– median
Kim *et al.* 2026	T2W	GYN hybrid	30	35 vol.	98.5 / –	0.74	2.52	37 s	mean
NeedleFinder 2026 (all 1 063)	3T T2W	GYN IC/IS	54	1063	91.4 / 86.9	0.47[j]	2.27[j]	**4.15 s**	median
NeedleFinder 2026 (759 long)	3T T2W	GYN IC/IS	54	759	91.3 / 88.7[k]	0.45[j]	2.41[j]	**4.15 s**	median

[a] Test-set catheters.
[b] Reported as Dice 0.88–0.90 and HD_{95} 0.71–0.77 mm.
[c] Root-mean-square distance to the clinical reconstruction.
[d] Fraction of catheters within 2 mm along 90 % of their length.
[e] Maximum Hausdorff distance below 2 mm for 92 % of catheters.
[f] Fraction within 3.4 mm.
[g] Mean Hausdorff distance.
[h] Cluster-level: 1.5 % false negatives and 6.7 % false positives among 826 catheters.
[i] Assisted reconstruction time, versus 46 min manual.
[j] Median (means: 0.59 ± 0.40 mm shaft and 6.23 ± 10.49 mm tip on all 1 063 catheters; 0.57 ± 0.39 mm shaft and 7.41 ± 12.23 mm tip on the 759 long shafts).
[k] Strict one-to-one recall at $\tau = 50$ % (92.5 % with inner-OOF cross-fit; 94.7 % at $\tau = 35$ %); precision counts predictions matched to short obturator tracks as correct (90.0 % $F1_{adj}$, 90.5 % with inner-OOF cross-fit). Times for **NeedleFinder 2026** are full-volume uncached per-patient averages across all 54 scans on CPU (4.15 s overall: 1.26 s on the 20 standard 320 × 256 scans, 5.85 s on the 34 high-resolution 640 × 416 / 768 × 768 scans) and one RTX 4090 GPU (2.67 s overall: 0.67 s standard, 3.85 s high-resolution).
[l] Treatments, not patients; 7207 of 24422 needles formed the evaluation set.
[m] Fully supervised model; their 10-shot synthetic-to-real variant reports 89.4 / 89.0, 0.33 mm and 0.90 mm.
[n] Signed along-axis tip displacement (mean ± SD 0.37 ± 1.68 mm), not an absolute 3D distance.

allows a 1D ResNet + self-attention localizer to pinpoint where the catheter void terminates and to quantify its own uncertainty—while incorporating sub-voxel inverse-variance-weighted polynomial fitting [55,62], directional axial TV continuity [43], and iterative multi-segment gap-bridging [59] refines shaft and tip accuracy and suppresses residual split fragments (4d, **89.4% F1**).

6.2 *How NeedleFinder 2026 Changes the Clinical Workflow in 3D Slicer*

Deployed as a one-click module in **3D Slicer** [82], NeedleFinder 2026 runs in **4.15 s** on a workstation CPU (1.26 s on standard-resolution scans) and **2.67 s** on a GPU. Instead of spending 30 to 45 minutes manually clicking 15 to 30 catheters slice by slice while the patient is under anesthesia, the medical physicist follows a structured, **under-one-minute simulated verification workflow**:

1. **Click *Detect Needles* (4 seconds):** The software crops the template plane automatically and reconstructs the 3D catheter bundle. Because our three-cornered-hat and non-inferiority analyses prove that the reconstructed shafts (0.39 mm intrinsic MAD SD vs. 0.35 mm and 0.48 mm for the two human expert teams) are statistically non-inferior to human inter-observer agreement, the shafts themselves rarely need manual editing.
2. **Check the 2 to 8 tips with the widest conformal uncertainty bars (16 to 64 seconds simulated):** The interface highlights the few catheter tips with large split-conformal uncertainty intervals ($\hat{\sigma}_{\text{tip}}$). Sliding the top 5 to 8 flagged endpoints along their already-traced shafts takes under a minute in simulation, drives 95th-percentile tip error below 10 mm (9.90 mm, mean 2.77 mm), and keeps counterfactual TG-43U1 target dose error $|\Delta D_{90}|$ to **2.8% median** (5.9% mean).
3. **Verify total catheter count (5 seconds simulated):** Since the clinician knows the exact number of needles inserted through the template (N_{implant}), if $N_{\text{pred}} < N_{\text{implant}}$, clicking the missed catheter (0.65 clicks per patient on average) raises 1-to-1 recall to **93.0%** (94.2% closest-match recall). Moreover, during progressive needle insertion in intraoperative MRI suites, the 1.31 s causal 4D Kalman mode can alert the implanting physician in real time if a newly inserted needle pushes a neighboring catheter out of position or deflects toward an organ at risk.

6.3 *Limitations and Future Directions*

Four limitations of the present study point directly to future work: 1. **Longitudinal tip ambiguity on thick slices.** While our 1D axial tip localizer cuts mean tip error to 6.23 mm and median tip error to 2.26–2.27 mm (58.7%–59.0% within ≤ 3.0 mm), 16.7% of unreviewed catheters still have a longitudinal tip error exceeding 10 mm when a catheter passes through a multi-slice signal dropout or terminates directly beneath dark bowel gas on 3 to 6 mm thick slices. Although split-conformal uncertainty triage (k = 5–8 tips reviewed) effectively contains this tail clinically, combining T2W planning scans with thin-slice isotropic 3D sequences (≤ 1 mm) or zero-echo-time sequences [13,18] is the natural path to eliminating unreviewed tip outliers altogether. 2. **Single-institution clinical implant cohort.** While we demonstrated

zero-shot transfer across 48 external public prostate scans, 3T prostate brachytherapy MRI, multiple MRI pulse sequences, 4D insertion series, and paired CT scans, all implanted brachytherapy patients were treated at Brigham and Women's Hospital. Prospective multi-center validation across 1.5T, 3T, and emerging low-field MRI systems [85] from multiple scanner vendors is the essential next step. 3. **Standardized TG-43U1 anatomical envelopes and simulated review timings.** Because clinical DICOM RTSTRUCT and RTPLAN files were not archived across the historical 54-patient repository, our TG-43U1 dosimetric study used standardized target and organ-at-risk envelopes derived from the active implant geometry, and our clinician tip-review curves used a quantitative 1-click oracle rather than prospective stopwatch trials in the operating room. 4. **Role of synthetic bundle pre-training (`SynthBrachy3T`).** Once Stage 1 is pre-trained on 51 external pelvic MRI scans and paired with collinear fragment stitching and the 1D tip localizer, `SynthBrachy3T` provides its largest detection boost when clinical bundle labels are scarce ($k = 0$ to 5 patients, lifting zero-shot long-shaft $F1_{\text{adj}}$ by +6.3 points to 83.4%), while offering a smaller refinement at the full 43-patient training budget (+0.4 F1 points, while reducing dwell-weighted tip error from 6.65 to 5.72 mm and counterfactual $|\Delta D_{90}|$ from 12.8% to 8.2%).

7 CONCLUSION

NeedleFinder 2026 reconstructs interstitial brachytherapy catheters directly on clinical 3T T2-weighted MRI by combining an analytical 3D dark-core symmetry prior (Stage 0), a self-supervised and multi-rater trained 3D residual U-Net (Stage 1), and a multi-stage graph and 1D axial tip localization pipeline enhanced with literature-inspired inverse-variance polynomial fitting, epistemic uncertainty continuity, and multi-segment gap-bridging (Stage 2). Evaluated under strictly nested 5 × 5 cross-validation across 54 gynecologic brachytherapy patients (1,063 catheters), the method achieves **89.1%–89.4% F1** (**90.0%–90.5%** $F1_{\text{adj}}$ on the 759 full-depth active-dwell shafts and **91.3%–91.8% F1** on the 23 patients with severe catheter adhesion), 0.47 **mm median** (0.58–0.59 **mm mean**) **in-plane shaft error**, and 2.26–2.27 **mm median** (6.23 **mm mean**) **3D distal tip error** in **4.15 s** on a workstation CPU (1.26 s on standard-resolution scans) and **2.67 s** on a GPU—significantly outperforming five peer-reviewed MRI/curvilinear methods and one recent tracking preprint reproduced on the identical splits. Across 40 doubly annotated clinical scans and 38 paired post-implant CT patients, three-cornered-hat analysis, one-sided non-inferiority tests, and CT adjudication demonstrate that NeedleFinder 2026 is statistically non-inferior to human expert agreement, while its split-conformal tip uncertainty intervals enable an under-one-minute simulated verification workflow in 3D Slicer that preserves AAPM TG-43U1 target dosimetry (2.8% median $|\Delta D_{90}|$) and generalizes zero-shot to prostate MRI, secondary MRI sequences, 4D intra-procedural insertion monitoring, and pelvic CT.

DATA AVAILABILITY

The 54 clinical MRI scans and their expert annotations were acquired at Brigham and Women's Hospital under Institutional Review Board approval and contain protected health information; they cannot be shared publicly. De-identified trajectory annotations and per-patient evaluation results may be made available by the corresponding author upon reasonable request, subject to institutional ethics approval. The 48 needle-free pelvic MRI scans are publicly available from the Medical Segmentation Decathlon (`Task05_Prostate`, http://medicaldecathlon.com [80]). The `SynthBrachy3T` synthetic bundle generator is released with our open-source repository.

CODE AVAILABILITY

The 3D Slicer extension, the native Rust Stage 0 kernel (`needlefinder_rust`), and the exported ONNX models (`NeedleUNet3D`, `NeedleSliceMPNN`, `TipAware BundleGraphTransformer`, and `Dedicated1DTipLocalizer`) are available at https://github.com/needlefinder/NeedleFinder under an open-source license. All scripts for reference-free preprocessing (`prepare_clean_dataset.py`), self-supervised pre-training (`pretrain_ssl_unet3d.py`), strictly nested 5×5 training and evaluation (`train_nested_stage1.py`, `run_stage2_and_tip_clean.py`, `export_clean_unified_and_table3.py`), peer-reviewed baseline reproduction and literature fusion (`run_part1_peer_baselines_and_unc_maps.py`, `run_part2_literature_fusion_inside_needlefinder.py`), specificity benchmarking (`benchmark_public_pelvic_mri_clean.py`), runtime profiling (`benchmark_54cases_timing_clean.py`), calibration and failure auditing (`calibration_and_failure_clean.py`), and clinical translation Experiments 1–4 (`run_exp1_tg43_dosimetry.py` through `run_exp4_4d_insertion_and_ct.py`) are archived in `vault/paper/benchmarks_clean/`.

ETHICS STATEMENT

This retrospective study analyzed imaging data acquired during routine clinical care under a protocol approved by the Institutional Review Board of Brigham and Women's Hospital / Mass General Brigham.

FUNDING

[Funding sources and grant numbers to be inserted.]

ACKNOWLEDGMENTS

The authors thank the clinical physics, radiation oncology, and radiology staff of the AMIGO suite at Brigham and Women's Hospital for their support in data acquisition and clinical annotation.

AUTHOR CONTRIBUTIONS

[To be completed following the CRediT taxonomy.]

CONFLICTS OF INTEREST

The authors declare no conflicts of interest.

REFERENCES

[1] Pötter R, Tanderup K, Kirisits C, Leeuw A de, Kirchheiner K, Nout R, Tan L T, Haie-Meder C, Mahantshetty U, Segedin B, et al 2018 The EMBRACE II study: The outcome and prospect of two decades of evolution within the GEC-ESTRO gynaecological working group and the EMBRACE studies *Clinical and Translational Radiation Oncology* **9** 48–60

[2] Pötter R, Tanderup K, Schmid M P, Jürgenliemk-Schulz I, Haie-Meder C, Fokdal L U, Sturdza A E, Hoskin P, Mahantshetty U, Segedin B, et al 2021 MRI-guided adaptive brachytherapy in locally advanced cervical cancer (EMBRACE-I): A multicentre prospective cohort study *The Lancet Oncology* **22** 538–47

[3] Kapur T, Egger J, Damato A, Schmidt E J and Viswanathan A N 2012 3-t MR-guided brachytherapy for gynecologic malignancies *Magnetic resonance imaging* **30** 1279–90

[4] Mastmeyer A, Pernelle G, Ma R, Barber L and Kapur T 2017 Accurate model-based segmentation of gynecologic brachytherapy catheter collections in MRI-images *Medical Image Analysis* **42** 173–88

[5] Zaffino P, Pernelle G, Mastmeyer A, Mehrtash A, Zhang H, Kikinis R, Kapur T and Spadea M F 2019 Fully automatic catheter segmentation in MRI with 3D convolutional neural networks: Application to MRI-guided gynecologic brachytherapy *Physics in Medicine & Biology* **64** 165008

[6] Chino J, Erickson B, Gaffney D, Hsu I-C, Kamrava M, Kim Y, Niedermayr T R, Roumeliotis M, Rownd J, Todor D, Viswanathan A and Kidd E A 2025 The American Brachytherapy Society (ABS) consensus guidance for hybrid intracavitary interstitial brachytherapy for locally advanced cervical cancer *Brachytherapy* **24** 463–78

[7] JC D, P P, K T, P P, D B, C K, EM P, Limbergen E van, C H-M and R P 2012 Recommendations from gynaecological (GYN) GEC-ESTRO working group (IV): Basic principles and parameters for MR imaging within the frame of image based adaptive cervix cancer brachytherapy *Radiotherapy and oncology : journal of the European Society for Therapeutic Radiology and Oncology* **103** 113–22

[8] Jacobsen M C, Beriwal S, Dyer B A, Klopp A H, Lee S I, McGinnis G J, Robbins J B, Rauch G M, Sadowski E A, Simiele S J, Stafford R J, Taunk N K, Yashar C M and Venkatesan A M 2022 Contemporary image-guided cervical cancer brachytherapy: Consensus imaging recommendations from the society of abdominal radiology and the american brachytherapy society *Brachytherapy* **21** 369–88

[9] Prisciandaro J, Zoberi J E, Cohen G, Kim Y, Johnson P, Paulson E, Song W, Hwang K-P, Erickson B, Beriwal S, Kirisits C and Mourtada F 2022 AAPM task group report 303 endorsed by the ABS: MRI implementation in HDR brachytherapy—considerations from simulation to treatment *Medical Physics* **49** e983–1023

[10] Richardson S L, Baghwala A, Cohen G N, Dempsey C, Libby B, Melhus C S, Miller R A, Scanderbeg D J and Simiele S J 2025 AAPM medical physics practice guideline 13.a: HDR brachytherapy, part B *Journal of Applied Clinical Medical Physics* **26** e70118

[11] Shaaer A, Paudel M, Smith M, Tonolete F, Nicolae A, Leung E and Ravi A 2020 Evaluation of an MR-only interstitial gynecologic brachytherapy workflow using MR-line marker for catheter reconstruction *Brachytherapy* **19** 642–50

[12] Malajovich I, Anamalayil S, Dolney O V, Kevin Teo B K, Arscott W T and Taunk N K 2020 Techniques for and uncertainties of MRI-based reconstruction of titanium tandem and ring brachytherapy applicators *Brachytherapy* **19** 651–8

[13] Lee C Y, Kaza E, Harris T C, O'Farrell D A, King M T, Dyer M A, Cormack R A and Buzurovic I 2023 Catheter reconstruction and dosimetric verification of MRI-only treatment planning (MRTP) for interstitial HDR brachytherapy using PETRA sequence *Physics in Medicine & Biology* **68** 035002

[14] Hellebust T P, Kirisits C, Berger D, Pérez-Calatayud J, De Brabandere M, De Leeuw A, Dumas I, Hudej R, Lowe G, Wills R and Tanderup K 2010 Recommendations from gynaecological (GYN) GEC-ESTRO working group: Considerations and pitfalls in commissioning and applicator reconstruction in 3D image-based treatment planning of cervix cancer brachytherapy *Radiotherapy and Oncology* **96** 153–60

[15] Ladbury C, Harkenrider M, Taunk N, Fisher C, Mayadev J, Venkat P, Yashar C, Gaffney D, Beriwal S and Glaser S 2023 A practical guide to hybrid interstitial/intracavitary brachytherapy for locally-advanced cervical cancer *Brachytherapy* **22** 640–8

[16] Zhang H, Donnelly E D, Strauss J B, Kang Z, Gopalakrishnan M, Lee P C, Khelashvili G, Nair C K, Lee B H and Sathiaseelan V 2019 Clinical implementation, logistics and workflow guide for MRI image based interstitial HDR brachytherapy for gynecological cancers *Journal of applied clinical medical physics* **20** 37–49

[17] Ning M S, Vedam S, Ma J, Stafford R J, Bruno T L, Cunningham M, Starks C, Lawyer A, Venkatesan A M, Wang L, Wang J, Olivieri N D, Guzman A B, Incalcaterra J R, Thaker N G, Joyner M M, Lin L L, Jhingran A, Eifel P J and Klopp A H 2020 Clinical utility and value contribution of an MRI-positive line marker for image-guided brachytherapy in gynecologic malignancies *Brachytherapy* **19** 305–15

[18] Kaza E, Lee C Y, King M T, Dyer M A, Cormack R A and Buzurovic I 2022 First pointwise encoding time reduction with radial acquisition (PETRA) implementation for catheter detection in interstitial high-dose-rate (HDR) brachytherapy *Brachytherapy* **21** 501–10

[19] Tanderup K, Hellebust T P, Lang S, Granfeldt J, Pötter R, Lindegaard J C and Kirisits C 2008 Consequences of random and systematic reconstruction uncertainties in 3D image based brachytherapy in cervical cancer *Radiotherapy and Oncology* **89** 156–63

[20] Schindel J, Zhang W, Bhatia S K, Sun W and Kim Y 2013 Dosimetric impacts of applicator displacements and applicator reconstruction-uncertainties on 3D image-guided brachytherapy for cervical cancer *Journal of Contemporary Brachytherapy* **5** 250–7

[21] Singh A, Karshafian R, Ravi A and Paudel M R 2026 Dosimetric evaluation of systematic source position uncertainty along the source path and robust planning during MRI-only cervical HDR brachytherapy *Journal of Applied Clinical Medical Physics* **27** e70810

[22] Rylander S, Buus S, Pedersen E M, Bentzen L and Tanderup K 2017 Dosimetric impact of contouring and needle reconstruction uncertainties in US-, CT- and MRI-based high-dose-rate prostate brachytherapy treatment planning *Radiotherapy and oncology : journal of the European Society for Therapeutic Radiology and Oncology* **123** 125–32

[23] Nunez D A, Trager M, Beaudry J, Cohen G N, Dauer L T, Gorovets D, Hassan Rezaeian N, Kollmeier M A, Leong B, McCann P, Williamson M, Zelefsky M J and Damato A L 2021 Quantifying clinical severity of physics errors in high-dose rate prostate brachytherapy using simulations *Brachytherapy* **20** 1062–9

[24] S H, SK N, JC L, J G and K T 2009 Applicator reconstruction in MRI 3D image-based dose planning of brachytherapy for cervical cancer *Radiotherapy and oncology : journal of the European Society for Therapeutic Radiology and Oncology* **91** 187–93

[25] Androulakis I, Godart J, Luthart L, MEMC C, Westerveld H, Nout R A, Hoogeman M and Kolkman-Deurloo I K 2025 Reconstruction errors in clinical intraoperative TRUS-based prostate HDR-BT detected using electromagnetic tracking *Brachytherapy* **24** 177–85

[26] Richart J, Carmona-Meseguer V, García-Martínez T, Herreros A, Otal A, Pellejero S, Tornero-López A and Pérez-Calatayud J 2018 Review of strategies for MRI based reconstruction of endocavitary and interstitial applicators in brachytherapy of cervical cancer *Reports of practical oncology and radiotherapy : journal of Greatpoland Cancer Center in Poznan and Polish Society of Radiation Oncology* **23** 547–61

[27] Shaaer A, Paudel M, Smith M, Tonolete F and Ravi A 2022 Deep-learning-assisted algorithm for catheter reconstruction during MR-only gynecological interstitial brachytherapy *Journal of Applied Clinical Medical Physics* **23** e13494

[28] Roumeliotis M, Antaki M, Kim G, Gardner U G, Sheikh K, Lee J, Seethamraju R, Tokuda J, Hu C, Schmidt E and Viswanathan A N 2026 Prospective operational feasibility, safety, and workflow of magnetic resonance-guided tracking technology for interstitial gynecologic brachytherapy *Brachytherapy* **25** 741–8

[29] Sargazi V, Naseri S, Gholamhosseinian H and Momennezhad M 2026 Applicator reconstruction in cervical cancer brachytherapy: A systematic review of current methods, challenges, and AI-driven future directions *Brachytherapy* **25** 262–74

[30] Vijayan S, Theeler J J, Fu Z and Kim Y 2026 Systematic review of artificial intelligence in brachytherapy *Journal of Contemporary Brachytherapy* **18** 85–114

[31] Chen X and Egger J 2014 Development of an open source software module for enhanced visualization during MR-guided interstitial gynecologic brachytherapy *SpringerPlus* **3** 167

[32] Pernelle G, Mehrtash A, Barber L, Damato A, Wang W, Seethamraju R T, Schmidt E J, Cormack R A, Wells W M, Viswanathan A N and Kapur T 2013 Validation of catheter segmentation for MR-guided gynecologic cancer brachytherapy *Medical image computing and computer-assisted intervention – MICCAI 2013* Lecture notes in computer science vol 8151 (Springer) pp 380–7

[33] Mastmeyer A, Pernelle G, Barber L, Pieper S, Fortmeier D, Wells S, Handels H and Kapur T 2017 Model-based catheter segmentation in MRI-images

[34] W Q, M Y, M D, D T and A F 2013 Needle segmentation using 3D hough transform in 3D TRUS guided prostate transperineal therapy *Medical physics* **40** 042902

[35] Dise J, Liang X, Scheuermann J, Anamalayil S, Mesina C, Lin L L and Teo B K 2015 Development and evaluation of an automatic interstitial catheter digitization tool for adaptive high-dose-rate brachytherapy *Brachytherapy* **14** 619–25

[36] J R, A O, S R, AI N, M D, M S, J V, F B and J P-C 2015 A practical MRI-based reconstruction method for a new endocavitary and interstitial gynaecological template *Journal of contemporary brachytherapy* **7** 407–14

[37] Hrinivich W T, Morcos M, Viswanathan A and Lee J 2019 Automatic tandem and ring reconstruction using MRI for cervical cancer brachytherapy *Medical Physics* **46** 4324–32

[38] Deufel C L, Tian S, Yan B B, Vaishnav B D, Haddock M G and Petersen I A 2020 Automated applicator digitization for high-dose-rate cervix brachytherapy using image thresholding and density-based clustering *Brachytherapy* **19** 111–8

[39] Zheng Y, Jiang S, Yang Z and Wei L 2021 Automatic needle detection using improved random sample consensus in CT image-guided lung interstitial brachytherapy *Journal of Applied Clinical Medical Physics* **22** 121–31

[40] Isensee F, Jaeger P F, Kohl S A A, Petersen J and Maier-Hein K H 2021 nnU-Net: A self-configuring method for deep learning-based biomedical image segmentation *Nature Methods* **18** 203–11

[41] Jung H, Shen C, Gonzalez Y, Albuquerque K and Jia X 2019 Deep-learning assisted automatic digitization of interstitial needles in 3D CT image based high dose-rate brachytherapy of gynecological cancer *Physics in Medicine & Biology* **64** 215003

[42] Jung H, Gonzalez Y, Shen C, Klages P, Albuquerque K and Jia X 2019 Deep-learning-assisted automatic digitization of applicators in 3D CT image-based high-dose-rate brachytherapy of gynecological cancer *Brachytherapy* **18** 841–51

[43] Dai X, Lei Y, Zhang Y, Qiu R L J, Wang T, Dresser S A, Curran W J, Patel P, Liu T and Yang X 2020 Automatic multi-catheter detection using deeply supervised convolutional neural network in MRI-guided high-dose-rate prostate brachytherapy *Medical Physics* **47** 4115–24

[44] Zhang D, Yang Z, Jiang S, Zhou Z, Meng M and Wang W 2020 Automatic segmentation and applicator reconstruction for CT-based brachytherapy of cervical cancer using 3D convolutional neural networks *Journal of Applied Clinical Medical Physics* **21** 158–69

[45] Zhang Y, He X, Tian Z, Jeong J J, Lei Y, Wang T, Zeng Q, Jani A B, Curran W J, Patel P, Liu T and Yang X 2020 Multi-needle detection in 3D ultrasound images using unsupervised order-graph regularized sparse dictionary learning *IEEE transactions on medical imaging* **39** 2302–15

[46] Zhang Y, Tian Z, Lei Y, Wang T, Patel P, Jani A B, Curran W J, Liu T and Yang X 2020 Automatic multi-needle localization in ultrasound images using large margin mask RCNN for ultrasound-guided prostate brachytherapy *Physics in Medicine & Biology* **65** 205003

[47] Zhang Y, Lei Y, Qiu R L J, Wang T, Wang H, Jani A B, Curran W J, Patel P, Liu T and Yang X 2020 Multi-needle localization with attention u-net in US-guided HDR prostate brachytherapy *Medical physics* **47** 2735–45

[48] Andersén C, Rydén T, Thunberg P and Lagerlöf J H 2020 Deep learning-based digitization of prostate brachytherapy needles in ultrasound images *Medical Physics* **47** 6414–20

[49] Wang F, Xing L, Bagshaw H, Buyyounouski M and Han B 2020 Deep learning applications in automatic needle segmentation in ultrasound-guided prostate brachytherapy *Medical Physics* **47** 3797–805

[50] Li X, Young A S, Raman S S, Lu D S, Lee Y H, Tsao T C and Wu H H 2020 Automatic needle tracking using mask r-CNN for MRI-guided percutaneous interventions *International journal of computer assisted radiology and surgery* **15** 1673–84

[51] Hu H, Yang Q, Li J, Wang P, Tang B, Wang X and Lang J 2021 Deep learning applications in automatic segmentation and reconstruction in CT-based cervix brachytherapy *Journal of Contemporary Brachytherapy* **13** 325–30

[52] Weishaupt L L, Sayed H K, Mao X, Choo R, Stish B J, Enger S A and Deufel C L 2022 Approaching automated applicator digitization from a new angle: Using sagittal images to improve deep learning accuracy and robustness in high-dose-rate prostate brachytherapy *Brachytherapy* **21** 520–31

[53] Liu D, Tupor S, Singh J, Chernoff T, Leong N, Sadikov E, Amjad A and Zilles S 2022 The challenges facing deep learning-based catheter localization for ultrasound guided high-dose-rate prostate brachytherapy *Medical Physics* **49** 2442–51

[54] Xie H, Wang J, Chen Y, Tu Y, Chen Y, Zhao Y, Zhou P, Wang S, Bai Z and Tang Q 2023 Automatic reconstruction of interstitial needles using CT images in post-operative cervical cancer brachytherapy based on deep learning *Journal of Contemporary Brachytherapy* **15** 134–40

[55] Aleong A M, Berlin A, Borg J, Helou J, Beiki-Ardakani A, Rink A, Raman S, Chung P and Weersink R A 2024 Rapid multi-catheter segmentation for magnetic resonance image-guided catheter-based interventions *Medical Physics* **51** 5361–73

[56] Wang Y, Jian W, Zhu L, Cai C, Zhang B and Wang X 2024 Attention-gated deep-learning-based automatic digitization of interstitial needles in high-dose-rate brachytherapy for cervical cancer *Advances in Radiation Oncology* **9** 101340

[57] Zhou W, Li X, Zabihollahy F, Lu D S and Wu H H 2024 Deep learning-based automatic pipeline for 3D needle localization on intra-procedural 3D MRI *International Journal of Computer Assisted Radiology and Surgery* **19** 2227–37

[58] Goulet M, Duguay-Drouin P, Mascolo-Fortin J, Mégrourèche J, Octave N and Tsui J M G 2025 Clinical application of deep learning-assisted needles reconstruction in prostate ultrasound brachytherapy *International Journal of Radiation Oncology, Biology, Physics* **122** 199–207

[59] Kostoulas V, Guijt A, Kerkhof E M, Pieters B R, Bosman P A N and Alderliesten T 2025 Dealing with segmentation errors in needle reconstruction for MRI-guided brachytherapy *Medical imaging 2025: Image-guided procedures, robotic interventions, and modeling, proc. SPIE 13408* p 1340826

[60] Moradi M M, Siavashpour Z, Takhtardeshir S, Showkatian E, Jaberi R, Ghaderi R, Mofid B and Taghizadeh-Hesary F 2025 Fully automatic reconstruction of prostate high-dose-rate brachytherapy interstitial needles using two-phase deep learning-based segmentation and object tracking algorithms *Clinical and Translational Radiation Oncology* **51** 100925

[61] Quetin S, Jafarzadeh H, Kalinowski J, Bekerat H, Bahoric B, Maleki F and Enger S A 2025 Automatic catheter digitization in breast brachytherapy *Medical Physics* **52** e18107

[62] Kim G, Antaki M, Schmidt E J, Roumeliotis M, Viswanathan A N and Lee J 2026 Automatic digitization of applicator and catheters for MRI-guided cervical cancer brachytherapy *Brachytherapy* **25** 85–94

[63] Hu X, Li F, Samaras D and Chen C 2019 Topology-preserving deep image segmentation *Advances in neural information processing systems 32 (NeurIPS 2019)* vol 32

[64] Shit S, Paetzold J C, Sekuboyina A, Ezhov I, Unger A, Zhylka A, Pluim J P W, Bauer U and Menze B H 2021 clDice – a novel topology-preserving loss function for tubular structure segmentation *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition (CVPR)* pp 16555–64

[65] Kirchhoff Y, Rokuss M R, Roy S, Kovacs B, Ulrich C, Wald T, Zenk M, Vollmuth P, Kleesiek J, Isensee F and Maier-Hein K 2024 Skeleton recall loss for connectivity conserving and resource efficient segmentation of thin tubular structures *Computer vision – ECCV 2024, lecture notes in computer science* (Springer) pp 218–34

[66] Xiao Z, Liu B, Wang J, Yao Q, Sun H, Zhou F, Wang J, Wu Q and Jiang P 2026 Multi-catheter digitization in brachytherapy via few-shot synthetic-to-real learning and structure-aware tracking *arXiv preprint arXiv:2607.11290*

[67] Xiao Z, Zhou F, Wang J, He C, Liu B, Sun H, Ji Z, Jiang Y, Wang J and Wu Q 2026 Multi-needle localization for pelvic seed implant brachytherapy based on tip-handle detection and matching *IEEE Journal of Biomedical and Health Informatics* 1–14

[68] Brasó G and Leal-Taixé L 2020 Learning a neural solver for multiple object tracking *Proceedings of the IEEE/CVF conference on computer vision and pattern recognition (CVPR)* pp 6246–56

[69] Ben-Haim T and Riklin Raviv T 2022 Graph neural network for cell tracking in microscopy videos *Computer vision – ECCV 2022* Lecture notes in computer science vol 13681 (Springer) pp 610–26

[70] Shit S, Koner R, Wittmann B, Paetzold J, Ezhov I, Li H, Pan J, Sharifzadeh S, Kaissis G, Tresp V and Menze B 2022 Relationformer: A unified framework for image-to-graph generation *Computer vision – ECCV 2022, lecture notes in computer science* (Springer) pp 422–39

[71] Zhou W, Dai Q, Curiel O, Tsao T C, Chiang J, Lu D S and Wu H H 2026 Keypoint detection network for needle localization on intra-procedural MRI in MRI-guided liver interventions *International journal of computer assisted radiology and surgery* **21** 1775–88

[72] Arapi V, Hardt-Stremayr A, Weiss S and Steinbrener J 2023 Bridging the simulation-to-real gap for AI-based needle and target detection in robot-assisted ultrasound-guided interventions *European Radiology Experimental* **7** 30

[73] Mehrtash A, Damato A, Pernelle G, Barber L, Farhat N, Viswanathan A, Cormack R and Kapur T 2014 EM-navigated catheter placement for gynecologic brachytherapy: An accuracy study *Medical imaging 2014: Image-guided procedures, robotic interventions, and modeling, proc. SPIE 9036* vol 9036 p 90361F

[74] Wang W, Viswanathan A N, Damato A L, Chen Y, Tse Z, Pan L, Tokuda J, Seethamraju R T, Dumoulin C L, Schmidt E J and Cormack R A 2015 Evaluation of an active magnetic resonance tracking system for interstitial brachytherapy *Medical physics* **42** 7114–21

[75] Arcos J de, Schmidt E J, Wang W, Tokuda J, Vij K, Seethamraju R T, Damato A L, Dumoulin C L, Cormack R A and Viswanathan A N 2017 Prospective clinical implementation of a novel magnetic resonance tracking device for real-time brachytherapy catheter positioning *International journal of radiation oncology, biology, physics* **99** 618–26

[76] Dürrbeck C, Sauer B C, Lotter M, Kreppner S, Strnad V, Fietkau R and Bert C 2022 Automating implant reconstruction in interstitial brachytherapy of the breast: A hybrid approach combining electromagnetic tracking and image segmentation *Radiotherapy and Oncology* **176** 172–8

[77] Cantin A, Lavallée M-C, Poulin E, Hautvast G, Foster W and Beaulieu L 2024 Proof-of-concept of real-time electromagnetic guidance for gynecologic interstitial catheters in high dose rate brachytherapy *Physics and Imaging in Radiation Oncology* **32** 100661

[78] Deufel C L, Brost E E, Dupere J M, Petersen I A, Haddock M G and Garda A E 2024 EMVision: An electromagnetic tracking guidance system for accurate placement of interstitial brachytherapy applicators *Brachytherapy* **23** 676–86

[79] Beaulieu L, Bert C, Borot de Battisti M, Cormack R A, Cunha J A M, Damato A L, Deufel C L, Hautvast G L T F, Hsu I-C, Kolkman-Deurloo I-K, Moerland M A, Niatsetski Y, Weersink R A and Tanderup K 2025 AAPM task group 317 report: A joint AAPM and ESTRO report on brachytherapy catheter, needle, and applicator tracking technology *Medical Physics* **52** e70037

[80] Antonelli M, Reinke A, Bakas S, Farahani K, Kopp-Schneider A, Landman B A, Litjens G, Menze B, Ronneberger O, Summers R M, et al 2022 The medical segmentation decathlon *Nature Communications* **13** 4128

[81] C K, MJ R, D B, F B, M D B, Laarse R van der, Y N, P P, TP H, J P-C, K T, JL V and FA S 2014 Review of clinical brachytherapy uncertainties: Analysis guidelines of GEC-ESTRO and the AAPM *Radiotherapy and oncology : journal of the European Society for Therapeutic Radiology and Oncology* **110** 199–212

[82] Fedorov A, Beichel R, Kalpathy-Cramer J, Finet J, Fillion-Robin J-C, Pujol S, Bauer C, Jennings D, Fennessy F, Sonka M, et al 2012 3D Slicer as an image computing platform for the quantitative imaging network *Magnetic Resonance Imaging* **30** 1323–41

[83] Kallis K, Kaltsas T, Kreppner S, Lotter M, Strnad V, Fietkau R and Bert C 2019 Impact of inter- and intra-observer variabilities of catheter reconstruction on multi-catheter interstitial brachytherapy of breast cancer patients *Radiotherapy and Oncology* **135** 25–32

[84] Karius A, Kreppner S, Strnad V, Schweizer C, Lotter M, Fietkau R and Bert C 2024 Inter-observer effects in needle reconstruction for temporary prostate brachytherapy: Dosimetric implications and adaptive CBCT-TRUS registration solutions *Brachytherapy* **23** 421–32

[85] Major T, Vízkeleti J and Takácsi-Nagy Z 2026 Geometric accuracy of applicator and needle reconstruction in gynecological brachytherapy using low-field magnetic resonance imaging *Journal of Contemporary Brachytherapy* **18** 288–95